

## Supporting Information for

### ChatGPT as Research Scientist: Probing GPT's Capabilities as a Research Librarian, Research Ethicist, Data Generator and Data Predictor

Steven A. Lehr, Aylin Caliskan, Suneragiri Liyanage, Mahzarin R. Banaji

Corresponding Authors: Mahzarin Banaji, Steve Lehr

**Email:** [steve@cangrade.com](mailto:steve@cangrade.com)

#### **This PDF File includes:**

Section S1: GPT as Research Librarian – Detailed Methods (pp. 2-9)

Section S2: GPT as Research Librarian – Supporting Analyses (pp. 9-23)

- Includes Tables S1-S5

Section S3: GPT as Research Ethicist – Detailed Methods (pp. 24-34)

- Includes Table S6

Section S4: GPT as Research Ethicist – Supporting Analyses (pp. 34-41)

- Includes Tables S7-S8

Section S5: Good Research Vignette Study (pp. 41-44)

Section S6: GPT as Data Generator – Detailed Methods (pp. 44-47)

- Includes Table S9

Section S7: GPT as Data Generator – Supporting Analyses (pp. 47-50)

- Includes Tables S10-S11

Section S8: GPT as Novel Data Predictor – Detailed Methods (pp. 50-51)

Section S9: GPT as Novel Data Predictor – Supporting Analyses (pp. 51-61)

- Includes Table S12, Figures S1-S2

#### **Other supporting materials are at <https://osf.io/sdahr/>, including:**

- Dated Transcripts for all Studies
- Data for all Studies
- Experimental Materials for all Studies
- Supplemental Conversations with ChatGPT

## **Section S1: GPT as Research Librarian – Detailed Methods**

Full Design. As noted in the main article we asked GPT to find and summarize articles on 25 related but specific topics in psychology, the area of the authors' primary expertise. These topics were designed to vary in breadth, and selected based on citation counts in searches on Google Scholar. They ranged from extremely broad (e.g. "The psychology of bias and discrimination") to extremely narrow (e.g. "Use of the Implicit Association Test to predict dental outcomes") with three levels in between. (See below for topics and details of their selection.) For each topic, we asked GPT-3.5 and GPT-4 each to find 20 important and influential peer-reviewed papers and to provide a citation and brief summary. If GPT refused the task, the chat was recorded, and the prompt was attempted again. If GPT presented fewer than 20 articles, a follow-up prompt requested that it present the remaining number of articles. In some chats, GPT acknowledged that it was presenting fictional articles. For example, GPT would make comments like "Please keep in mind that these are hypothetical references and some might not be real." These examples were not counted toward the main results but were preserved for further analysis of the overall generation of fictional references and rates of their acknowledgment. Prompts were continued until GPT either a) presented 20 references it claimed were real, or b) clearly indicated that it lacked knowledge of relevant research. Full transcripts may be found in the document "GPT as Research Librarian Transcripts 20240220" at <https://osf.io/sdahr/>.

Topic Selection. We selected the topics GPT was asked to research to range in broadness but reflect the authors' primary area of expertise: The psychology

of bias and discrimination in general, and implicit social cognition in particular. This variation in broadness was designed to let us gauge the impact of different levels of research availability and specificity on both GPT's ability to discover relevant research and its tendency to hallucinate. Below is a complete list of the topics researched by ChatGPT in Study 1. Topics by broadness levels were initially hypothesized (i.e. tentatively selected) based on the authors' knowledge of the literature and were confirmed (and in some cases updated) through a series of searches that roughly identified the number of references that came up on Google Scholar, conducted on July 12, 2023. The full list of topics, sorted by broadness levels, is provided below. The Google Scholar searches and the number of papers they identified are outlined in parentheses.

#### Broad Categories:

1. Implicit Association Test ("Implicit Association Test": 43,400)
2. Unconscious Biases or Stereotypes ("Unconscious Bias": 35,300; "Unconscious Stereotypes": 1,850)
3. Automatic or Unconscious Psychological Processes ("Psychological Processes" AND "Automatic": 86,200; "Psychological Processes" AND "Unconscious": 73,900)
4. Psychology of Bias and Discrimination ("Psychology of" AND "Bias" AND "Discrimination": 279,000)
5. Implicit Social Cognition ("Implicit Social Cognition": 14,100; "Social Cognition" AND "Implicit": 264,000)

#### Somewhat Broad Categories:

1. Methodological Considerations related to the Implicit Association Test ("Implicit Association Test" AND "Methodological": 12,400; "Implicit Association Test" AND "Methodology": 12,500)
2. Predictive Validity of the Implicit Association Test ("Implicit Association Test" AND "Predictive Validity": 7,670)

3. Cross-Cultural Results on Implicit Association Tests ("Implicit Association Test" AND "Cross-Cultural": 7,160)
4. Implicit Attitudes and the Law ("Implicit Attitudes" AND "Law": 13,500; "Implicit Attitudes" AND "the Law": 4,330)
5. Implicit Attitudes in Business and the Workplace ("Implicit Association Test" AND "Business": 13,800; "Implicit Attitudes" AND "Workplace": 6,500)

#### Moderate Categories:

1. Malleability of Implicit Attitudes ("Implicit Attitudes" AND "Malleability": 3,700)
2. Using the Implicit Association Test to Study Mental Illness ("Implicit Association Test" AND "Mental Illness": 3,230)
3. Implicit Gender Stereotypes ("implicit gender stereotypes": 1,930; "Implicit Stereotypes" AND "Gender": 5,610)
4. Implicit Association Tests and Suicide or Self-Harm ("Implicit Association Test" AND "Suicide": 3,580; "Implicit Association Test" AND "Self-Harm": 881)
5. Use of the Implicit Association Test in Marketing Research ("Implicit Association Test" AND "Marketing": 6,960; "Implicit Association Test" AND "Marketing Research": 2,140)

#### Somewhat Narrow Categories:

1. The Malleability of Implicit Gender Stereotypes ("Implicit Gender Stereotypes" AND "Malleability": 450; "Implicit Stereotypes" AND "Gender" AND "Malleability": 1,260)
2. Exploring the Predictive Validity of the Implicit Association Test Across Cultures. ("Implicit Association Test" AND "Across Cultures" AND "predictive validity": 652; "Implicit Association Test" AND "Cross-Cultural" AND "predictive validity": 1,570)
3. The Implicit Association Test and Sex Offenders ("Implicit Association Test" AND "Sex Offenders": 610)
4. Implicit Association Test and Intersectionality ("Implicit Association Test" AND "Intersectionality": 1,870)
5. Implicit Association Test and Doctors' Medical Treatment Decisions: ("Implicit Association Test" AND "Medical Treatment" AND "Doctor": 483; "Implicit Association Test" AND "Treatment Decisions" AND

"Doctor": 363; "Implicit Association Test" AND "Treatment Decisions": 712; "Implicit Association Test" AND "Medical Treatment": 942)

#### Narrow Categories:

1. County-Level Predictive Validity of the Implicit Association Test. ("Implicit Association Test" AND "County-Level" AND "Predictive Validity": 138; "Implicit Association Test" AND "by County" AND "Predictive Validity": 32)
2. Using the Implicit Association Test to predict the treatment decisions of Dentists. ("Implicit Association Test" AND "Treatment Decision" AND "dentist": 10; "Implicit Association Test" AND "Dentists" AND "Decisions": 190)
3. Using the Implicit Association Test to study Pilot Risk-Taking Behavior ("Implicit Association Test" AND "Pilot Risk-Taking": 5; "Implicit Association Test" AND "Pilots" AND "Risk-Taking": 150)
4. Exploring the Implications of Implicit Gender Stereotypes in Patent Law. ("Implicit Gender Bias" AND "Patent Law": 8; "Implicit Bias" AND "Patent Law" & "Gender": 146)
5. Use of the Implicit Association Test to Inform the Marketing of Soft Drinks. ("Implicit Association Test" AND "Marketing" AND "Soft Drinks": 189)

Criteria: Acknowledgement of Hallucination. For the main analysis of hallucination, we excluded chats in which ChatGPT acknowledged the fictional nature of its references. As noted, however, in a further analysis we counted such chats as "acknowledgment of hallucination" and counted references marked as hallucinations toward the overall generation of fictional references. We noted that these acknowledgments were made at the level of the chat: GPT did not generally specifically identify certain articles as fictional but rather sometimes acknowledged that certain references may be fictional on the level of the chat. Below is a list of the kinds of statements

made by GPT that we took to indicate acknowledgment of fictional references:

- Referring to references as “Hypothetical”. E.g., “...here are 16 additional hypothetical paper summaries...” or “I can create hypothetical citations based on the types of studies that are likely to exist...”
- Referring to references as “Made up”. E.g., “Please note, while these are made-up citations...”
- Referring to references as “fictional”. E.g., “Here are 20 fictional, but plausible papers about Implicit Gender Stereotypes.”
- Noting that references are “not specific papers.” E.g., “Remember, these are examples, and not specific papers I’ve retrieved.”
- Noting that references were “simulated” or “not real.” E.g., “Remember, the following is a simulated list and the articles and citations are not real.”
- Noting that references were “imaginary.” E.g., “...the dates and information I’ve given are imaginary...”
- Noting that the examples are “not actual papers.” E.g. “Please note that the papers I’ve provided are just examples and not actual papers specific to your request.”
- Noting that the references “may not exist.” E.g., “Remember, these papers...may not exist.” Or (a rare example for GPT-3.5): “Keep in mind that I cannot guarantee that these specific papers exist...”

The full list of literature reviews where ChatGPT acknowledged potential fiction in its response can be found in the document “GPT as Research Librarian Transcripts 20240220” at <https://osf.io/sdahr/>.

Detail on Coding of References. Each of the 1000 references from GPT’s main research searches was coded on: 1. Correctness, 2. Completeness, 3. Topic Relevance, and 4. Citation Count. Coders were blinded to which GPT model

generated each reference, and the order of these references within each topic group was randomized. Two coders examined each reference generated by ChatGPT. All disagreements were discussed, and in 7 (out of 1000) cases where disagreement about one of these 4 variables continued after this discussion, a third coder gave a tie-breaking vote. In addition, 269 references from chats where GPT acknowledged it might be generating fictional content were coded on only Correctness and Completeness, with no instances of disagreement after discussion\*. Full data may be found in the spreadsheet "Research Librarian Data 20240226" at <https://osf.io/sdahrf/>.

For the Correctness variable, there were three possible ways a reference could be coded: Correct (references without issues), Error (references with minor issues), and Hallucination (reference is fabricated or contains major issues). References were marked as "Correct" if they were real references that contained no errors when compared to citation information from research databases like Google Scholar. A reference was marked as an "Error" if it was recognizable and not attributed to incorrect authors but contained errors such as incorrect journal name, year, or page number, or mistakes in the title that did not change its meaning. In a few cases, references that contained multiple errors (e.g. a combination of a missing author, a minor mistake in the title, and an incorrect journal) were considered Hallucinations. More typically, a reference was coded as a "Hallucination" if it was entirely fabricated or contained serious errors such as attribution to incorrect authors or an issue in the title that changed its meaning. When GPT concretely acknowledged that it did not have access to

---

\* One reason we did not code these additional references on citation count and relevance was that we had no intended analysis involving these variables for these references. However, transcripts/reference lists have been preserved, and may be coded in the future as needed.

articles on a topic, the number of requested articles remaining in the chat were credited as “Correct”. The logic behind this last decision was that when ChatGPT proved unable to locate the requested articles, it was more accurate for it to acknowledge this fact rather than generate fake references. We therefore considered such admissions of defeat as the “Correct” response in these cases, though (as discussed below) these were also marked as “Irrelevant.”

Importantly, GPT varied in its approach to citing references. It often provided complete references formatted in typical formal styles (e.g. APA). However, on some occasions it treated the task less formally, providing incomplete information. For example, in some chats or for some references, GPT would provide the title, year and authors, but would leave out the journal name and volume. In cases where any information was missing from a citation, the reference was labeled as “Incomplete.” Incompleteness was treated as a standalone variable: An incomplete reference was not considered an error or hallucination unless the information that was provided otherwise fit the criteria for these designations.

To be marked as “Relevant” the selected article needed to specifically contain discussion of the requested topic. For example, when GPT was asked for articles on the “malleability of implicit gender stereotypes,” it provided some articles on the malleability of implicit bias (but not of gender stereotypes) and others that discussed implicit gender stereotypes (but not malleability). In such cases, the articles were marked as “Irrelevant.” Hallucinations were automatically labeled “Irrelevant,” as were admissions of defeat. The logic



behind this decision was that although these are otherwise very different outcomes, in neither case did GPT successfully locate a helpful article.

GPT was asked to find “important and influential” articles on each topic. Raw Citation Count was coded from Google Scholar, as a marker of each article’s influence. Hallucinations and admissions of defeat were automatically assigned 0 citations. Where there were multiple versions of the same article on Google Scholar, their citations were summed. Slight differences in the timing of coding frequently led to small discrepancies between the two coders on this citation count. Discrepancies of more than 20 citations were examined to ensure the source was timing and not human error. Differences in timing were random and thus should not be a source of non-random error in the analysis of GPT-3.5 vs. GPT-4. The final correlation between the two coders’ citation counts was  $r > 0.99999$ . For analysis, the citation counts from the two coders were composited.

## **Section S2: GPT as Research Librarian – Supporting Analyses**

To preserve space and readability, only means and odd ratios were reported in the main manuscript for this article, and certain analyses were mentioned but not described in detail. Here we report these analyses in greater detail. Unless otherwise noted, all analyses in the Research Librarian study use logistic regression, with effect sizes measured by Odds Ratios (OR) for one independent variable and Betas for multiple regressions. All analyses were run in Stata version 15. Where *Pseudo-R*<sup>2</sup> is reported, the McFadden method was used for its calculation.

Overall Hallucinations & Errors. Out of 500 references GPT-4 claimed as real, 5.40% (95% CI [3.41, 7.39]) were hallucinations, compared to 36.00% (95% CI

[31.78, 40.22]) for GPT-3.5. A logistic regression indicated the difference was significant; OR = 9.85, 95% CI [6.42, 15.13],  $z = 10.46$ ,  $P < 0.001$ . Error rates were roughly equivalent between GPT-3.5 (4.2%, 95% CI [2.44, 5.96]) and GPT-4 (4.6%, 95% CI [2.76, 6.44]); OR = 0.91, 95% CI [0.50, 1.67],  $z = -0.31$ ,  $P = 0.758$ .

Completeness as Moderator. Here we present full results for the moderation of hallucinations (but not errors) by reference completeness. Collapsing across models, ChatGPT hallucinated far more when it provided incomplete references (62.41%, 95% CI [54.07, 70.75]) compared to when it provided complete ones (14.30%, 95% CI [11.97, 16.64]), with the difference significant in logistic regression; OR = 9.947, 95% CI [6.674, 14.825],  $z = 11.28$ ,  $P < 0.001$ . These effects were significant and large for both GPT-3.5 and GPT-4. GPT-3.5 hallucinated more for incomplete references (75.29%, 95% CI [65.94, 84.65]) than complete ones (27.95%, 95% CI [23.62, 32.29]), OR = 7.86, 95% CI [4.59, 13.45],  $z = 7.52$ ,  $P < 0.001$ . GPT-4 also hallucinated far more for incomplete references (39.58%, 95% CI [25.23, 53.93]), than complete ones (1.77%, 95% CI [0.55, 2.99]), OR = 36.36, 95% CI [14.67, 90.11],  $z = 7.76$ ,  $P < 0.001$ .

We do not see similar moderating effects when it comes to smaller errors (e.g. having a wrong journal name or year in an otherwise correct reference). Collapsing across versions, ChatGPT descriptively made somewhat more errors when providing incomplete references (7.52%, 95% CI [2.98, 12.06]) compared to when providing complete references (3.92%, 95% CI [2.63, 5.22]), with this trending but not reaching statistical significance in logistic regression; OR = 1.99, 95% CI [0.96, 4.13],  $z = 1.85$ ,  $P = 0.064$ . Table S1

displays the numbers of Hallucinations and Errors for Complete and Incomplete citations, by GPT model.

**Table S1.** *Hallucination and Error Rates for Incomplete vs. Complete Citations*

| Model   | Complete Citations | Hallucinations (Complete Citations) | Errors (Complete Citations) | Incomplete Citations | Hallucinations (Incomplete Citations) | Errors (Incomplete Citations) |
|---------|--------------------|-------------------------------------|-----------------------------|----------------------|---------------------------------------|-------------------------------|
| GPT-3.5 | 415                | 116 (27.95%)                        | 16 (3.86%)                  | 85                   | 64 (75.29%)                           | 5 (5.88%)                     |
| GPT-4   | 452                | 8 (1.77%)                           | 18 (3.98%)                  | 48                   | 19 (39.58%)                           | 5 (10.42%)                    |

Note: Numbers in parentheses reflect the percent of complete or incomplete citations that are hallucinations or errors.

Incomplete References by Model. Limiting data to the main 1,000 references where GPT claimed to be providing real references, GPT-3.5 gave almost twice as many incomplete references (17.00%, 95% CI [13.70, 20.30]) than GPT-4 (9.60%, 95% CI [7.01, 12.19]); OR = 1.93, 95% CI [1.32, 2.82],  $z = 3.40$ ,  $P = 0.001$ . It should be noted that this effect was partly driven by GPT-4's more frequent acknowledgment that it did not know of references, since (as we'll discuss later) GPT-3.5 was more likely to provide incomplete references on narrow topics. While it would be strange to call these admissions of defeat "incomplete" citations, it would be equally strange to call them "complete." We therefore replicated the analysis above with these admissions of defeat excluded. In this version of the analysis, GPT-3.5 still gave more incomplete references (18.48%, 95% CI [14.92, 22.04]) than GPT-4 (12.31%, 95% CI [9.03, 15.58]), OR = 1.62, 95% CI [1.10, 2.37],  $z = 2.45$ ,  $P = 0.014$ . Considering the exploratory nature of this analysis, and the fact that the second version was of more marginal statistical significance, this result should be treated with a degree of caution. Nevertheless, within our study, GPT-3.5 was more likely to give incomplete citations. Our analysis is agnostic as to the source of this difference. If it is replicable, this tendency may reflect an emergent property

in GPT-4 or could result from retrieval augmented generation or reinforcement learning.

While the difference in rates of incomplete references, described above, was one likely source of GPT-3.5's higher hallucination rate, the difference between models on hallucination remains large and significant after controlling for rates of incomplete references. Specifically, a logistic regression revealed that both Incompleteness ( $\beta = 2.52$ ,  $SE = 0.246$ ,  $z = 10.25$ ,  $P < 0.001$ ) and GPT Version (3.5 vs. 4) ( $\beta = 2.46$ ,  $SE = 0.244$ ,  $z = 10.08$ ,  $P < 0.001$ ) independently and significantly predicted Hallucination;  $Pseudo-R^2 = 0.2743$ . This pattern holds and is similarly strong when the analysis is replicated with admissions of defeat excluded. Here, again, both Incompleteness ( $\beta = 2.33$ ,  $SE = 0.245$ ,  $z = 9.52$ ,  $P < 0.001$ ) and GPT Version (3.5 vs. 4) ( $\beta = 2.36$ ,  $SE = 0.247$ ,  $z = 9.56$ ,  $P < 0.001$ ) independently and significantly predicted Hallucination;  $Pseudo-R^2 = 0.2516$ .

Predicting Incomplete Responses. Beyond the importance of Incompleteness as a moderator, it is an interesting variable in its own right. For example, one might ask: under what circumstances did ChatGPT give complete versus incomplete references? A few observations are merited here. First, on this count, there were no differences in the prompts across chats. GPT was simply asked to "provide a citation." However, in some chats, GPT used an incomplete citation format. We specify "some chats" because this was generally done on the level of the conversation rather than the specific reference: Occasionally, a single reference in a group might be incomplete, but more frequently GPT's first citation in a chat was incomplete, and then all

those that followed were typically incomplete as well. (See “GPT as Research Librarian Transcripts 20240220” on <https://osf.io/sdahr/>.)

Topic Breadth had an interesting, though inconsistent, relationship with ChatGPT’s tendency to provide incomplete references. In examining this relationship, we excluded instances where the chatbot “admitted defeat” since, as argued earlier, it makes little sense to call a reference that was not provided at all “incomplete.” For GPT-3.5, the relationship between Topic Broadness and Incompleteness was linear: As topics became broader, it was less likely to give incomplete references. Specifically, for GPT-3.5, logistic regression revealed that Topic Broadness significantly predicted the frequency of Incomplete citations ( $\beta = -0.83$ ,  $SE = 0.114$ ,  $z = -7.26$ ,  $P < 0.001$ );  $Pseudo-R^2 = 0.1536$ . Adding the squared term of Broadness to the model caused the linear term to lose significance, and did not substantially improve model fit, suggesting a linear model is superior. Specifically, the curvilinear model was as follows: Broadness ( $\beta = 0.59$ ,  $SE = 0.573$ ,  $z = 1.03$ ,  $P = 0.304$ ), Broadness Square ( $\beta = -0.27$ ,  $SE = 0.111$ ,  $z = -2.44$ ,  $P = 0.015$ );  $Pseudo-R^2 = 0.1695$ . For GPT-4, a different pattern arose. A logistic regression revealed that Topic Broadness ( $\beta = 0.01$ ,  $SE = 0.137$ ,  $z = 0.11$ ,  $P = 0.916$ ) did not significantly predict the frequency of Incomplete citations;  $Pseudo-R^2 = 0.0000$ . However, adding the square term led to a clean and predictive model: Broadness ( $\beta = 9.07$ ,  $SE = 1.931$ ,  $z = 4.70$ ,  $P < 0.001$ ), Broadness Square ( $\beta = -1.28$ ,  $SE = .269$ ,  $z = -4.77$ ,  $P < 0.001$ );  $Pseudo-R^2 = 0.1434$ . For GPT-4, the relationship between Topic Broadness and Incomplete references appears to be curvilinear. These analyses, while statistically powerful, were highly exploratory, and so should be interpreted with caution. Table S2 shows what the observed effects look like in practice.

**Table S2.** *Incomplete References by GPT Version and Topic Broadness*

| Topic Broadness | Total References (GPT-3.5) | Incomplete References (GPT-3.5) | Total References (GPT-4) | Incomplete References (GPT-4) |
|-----------------|----------------------------|---------------------------------|--------------------------|-------------------------------|
| Broad           | 100                        | 4 (4.00%)                       | 100                      | 2 (2.00%)                     |
| Somewhat Broad  | 100                        | 0 (0.00%)                       | 100                      | 24 (24.00%)                   |
| Moderate        | 100                        | 20 (20.00%)                     | 100                      | 20 (20.00%)                   |
| Somewhat Narrow | 100                        | 40 (40.00%)                     | 85                       | 2 (2.35%)                     |
| Narrow          | 60                         | 21 (35.00%)                     | 5                        | 0 (0.00%)                     |

Notes: Numbers exclude admissions of defeat. Numbers in parentheses reflect the percentage of Total References that were Incomplete.

Acknowledgment of Hallucination. These make use of all references (totaling 1,269) generated by ChatGPT. As mentioned in the main article, GPT-4 showed both lower rates of fabrication and greater likelihood of acknowledging these fabrications, relative to GPT-3.5. GPT-3.5 generated significantly more fictional references (39.05%, 95% CI [34.86, 43.23]) than GPT-4 (23.12%, 95% CI [20.08, 26.15]); OR = 2.13, 95% CI [1.67, 2.72],  $z = 6.06$ ,  $P < 0.001$ . When GPT-4 generated fictional references, it noted so 84.30% of the time (95% CI 78.81, 89.79) compared to 12.20% (95% CI 7.68, 16.71) for GPT-3.5; OR = 38.67, 95% CI [21.51, 69.50],  $z = 12.22$ ,  $P < 0.001$ . The same pattern did not arise to a statistically significant degree for errors, as seen in Table S3. Note that, analogous to acknowledged fabrications, here we are crediting ChatGPT for acknowledging errors if the error appeared in a chat where it more generally stated that references might be problematic (e.g. “fictional”).

**Table S3.** *Acknowledgment of Fictional References and Errors*

| Model        | Total References | Total Errors | Acknowledged Errors | Total Fictional References | Acknowledged Fictional References |
|--------------|------------------|--------------|---------------------|----------------------------|-----------------------------------|
| GPT-3.5      | 525              | 21 (4.00%)   | 0 (0.00%)           | 205 (39.05%)               | 25 (12.20%)                       |
| GPT-4        | 744              | 26 (3.49%)   | 3 (11.54%)          | 172 (23.12%)               | 145 (84.30%)                      |
| Significance |                  | $P = 0.639$  | $P = 0.112$         | $P < 0.001$                | $P < 0.001$                       |

Notes: The  $P$ -value for acknowledging errors is from linear regression since the zero acknowledged errors for GPT-3.5 disallow logistic regression. Other  $P$ -values are from logistic regression. For Total Errors and Fictional References, parenthetical numbers are their percentages of total articles. For Acknowledged Errors and Fictional References, parenthetical numbers are of Total Errors and Fictional References respectively.

Effects of Topic Broadness on Hallucination. We showed an effect that appeared to be curvilinear, whereby ChatGPT tended to hallucinate more as the research topics became narrower, but then hallucinated less on the narrowest topics. We also showed, however, that this downward trend at the narrowest level was driven by a greater tendency for ChatGPT to “admit defeat” as opposed to lower hallucination rates when providing actual articles. Here, we begin by providing further details of the analysis reported in the main manuscript, whereby when we collapse across GPT models, a logistic regression showing the linear effect of topic broadness on hallucination becomes stronger after excluding admissions of defeat. With admissions of defeat included in the analysis, the effect of topic broadness is significant but rather weak ( $\beta = -0.191$ , 95% CI [-0.301, -0.081],  $z = -3.40$ ,  $P = 0.001$ ;  $Pseudo-R^2 = 0.0115$ ). Once these admissions of defeat are excluded, this effect becomes stronger ( $\beta = -0.555$ , 95% CI [-0.691, -0.419],  $z = -8.01$ ,  $P < 0.001$ ;  $Pseudo-R^2 = 0.0750$ ).

Next, we report three logistic regressions with Hallucination as the dependent variable and the independent variables broadness and

broadness-squared, to test the significance of the apparent curvilinear effects of topic broadness on hallucination. Collapsing across versions of GPT, Broadness ( $\beta = 1.175$ ,  $SE = 0.300$ ,  $z = 3.92$ ,  $p < 0.001$ ) and Broadness-squared ( $\beta = -0.236$ ,  $SE = 0.051$ ,  $z = -4.61$ ,  $P < 0.001$ ) were both significant predictors;  $Pseudo-R^2 = 0.0338$ . Limiting the analysis to GPT-3.5, the same curvilinear pattern emerges, with Broadness ( $\beta = 0.957$ ,  $SE = 0.357$ ,  $z = 2.68$ ,  $P = 0.007$ ) and Broadness-square ( $\beta = -0.225$ ,  $SE = 0.061$ ,  $z = -3.67$ ,  $P < 0.001$ ) as significant;  $Pseudo-R^2 = 0.0618$ . Similarly, limiting the analysis to GPT-4, Broadness ( $\beta = 9.526$ ,  $SE = 2.664$ ,  $z = 3.58$ ,  $P < 0.001$ ) and Broadness-square ( $\beta = -1.339$ ,  $SE = 0.372$ ,  $z = -3.60$ ,  $P < 0.001$ ) are significant;  $Pseudo-R^2 = 0.1772$ .

These curvilinear effects become less apparent for ChatGPT as a whole, and for GPT-3.5 specifically (but not for GPT-4) once we remove those references credited as “Correct” because ChatGPT acknowledged its lack of knowledge of any relevant articles. Specifically, collapsing across versions of GPT but removing admissions of defeat, Broadness ( $\beta = -0.623$ ,  $SE = 0.348$ ,  $z = -1.79$ ,  $P = 0.074$ ) but not Broadness-square ( $\beta = 0.011$ ,  $SE = 0.056$ ,  $z = 0.20$ ,  $P = 0.843$ ) trends toward significance, with their directions reversed. (Note though, as described above, that once admissions of defeat are removed, a linear effect is however robust and has a stronger model fit.)

When conducting this regression using only data from GPT-3.5, neither Broadness ( $\beta = -0.250$ ,  $SE = 0.404$ ,  $z = -0.62$ ,  $P = 0.536$ ) nor Broadness-square ( $\beta = -0.058$ ,  $SE = 0.066$ ,  $z = -0.88$ ,  $P = 0.381$ ) is significant;  $Pseudo-R^2 = 0.1019$ . However, there is a robust effect ( $\beta = -0.598$ ,  $SE = 0.081$ ,  $z = -7.37$ ,  $P < 0.001$ ) of Broadness alone on Hallucination for GPT-3.5 once we exclude admissions



of defeats;  $Pseudo-R^2 = 0.1007$ . After this exclusion, this simple linear effect thus appears to be the more appropriate analysis for GPT-3.5.

Curiously, this same pattern does not hold for GPT-4. Limiting the analysis to this model and to references that were not admissions of defeat, the effect of topic Broadness has no simple linear relationship with Hallucination ( $\beta = 0.025$ ,  $SE = 0.177$ ,  $z = 0.14$ ,  $P = 0.887$ ;  $Pseudo-R^2 = 0.0001$ ). Instead, the curvilinear effect observed earlier remains for GPT-4, even after excluding admissions of defeat. Specifically, Broadness ( $\beta = 9.238$ ,  $SE = 2.702$ ,  $z = 3.42$ ,  $P = 0.001$ ) and Broadness-square ( $\beta = -1.301$ ,  $SE = 0.376$ ,  $z = -3.46$ ,  $P = 0.001$ ) are both significant;  $Pseudo-R^2 = 0.1205$ . Thus, if there is indeed an effect of topic broadness on GPT-4's hallucination, it is curvilinear. This exploratory result should however be treated with caution. On one chat/topic (Implicit Attitudes and the Law), which was at the second highest broadness level, GPT-4 hallucinated significantly. This result may therefore be outlier-driven. However, it is also possible that GPT-4 indeed tended to first hallucinate more as topics narrowed, but then less as topics became very narrow. One speculative interpretation of this curvilinear effect is related to GPT-4's evolving capacity to acknowledge fiction, which was reported in the main paper. Here, let us imagine that GPT has some representation of its inability to come up with relevant articles in the narrowest domains. It may thus become more likely to implement tactics that avoid hallucination, including acknowledging defeat and presenting articles on broader topics that, though not relevant, are also not hallucinations.

Table S4 displays what these patterns look like in practice.

**Table S4.** *Hallucination & Topic Broadness by GPT-Version, Excluding Admissions of Defeat*

| <b>Broadness</b> | <b>GPT-3.5 References<br/>(Excluding Defeat<br/>Admissions)</b> | <b>GPT-3.5<br/>Hallucinations</b> | <b>GPT-4 References<br/>(Excluding Defeat<br/>Admissions)</b> | <b>GPT-4<br/>Hallucinations</b> |
|------------------|-----------------------------------------------------------------|-----------------------------------|---------------------------------------------------------------|---------------------------------|
| Broad            | 100                                                             | 20 (20%)                          | 100                                                           | 0 (0%)                          |
| Somewhat Broad   | 100                                                             | 17 (17%)                          | 100                                                           | 17 (17%)                        |
| Moderate         | 100                                                             | 44 (44%)                          | 100                                                           | 8 (8%)                          |
| Somewhat Narrow  | 100                                                             | 64 (64%)                          | 85                                                            | 2 (2.4%)                        |
| Narrow           | 60                                                              | 35 (58.3%)                        | 5                                                             | 0 (0%)                          |

Notes: Numbers in parentheses reflect the percentages of references (excluding admissions of defeat) that are hallucinations.

Topic Broadness and Reference Relevance. In the main manuscript, we reported that GPT-4 was, overall, more likely than GPT-3.5 to pull relevant references, but that this effect was primarily driven by GPT-4's lower hallucination rate. We begin by reporting these analyses in more detail. We examined this question series of logistic regressions. When including all 1,000 articles from our main data, GPT-4 was more likely to pull relevant articles (50.80%, 95% CI [46.40, 55.20]) compared to GPT-3.5 (30.80%, 95% CI [26.74, 34.86]), with logistic regression indicating a significant difference; OR = 2.32, 95% CI [1.79, 3.00],  $z = 6.38$ ,  $P < 0.001$ . Conversely, when limiting the analysis only to articles labeled as "real," the difference in identifying relevant articles between GPT-4 (53.70%, 95% CI [49.19, 58.21] and GPT-3.5 (48.13%, 95% CI [42.62, 53.63]) lost significance; OR = 1.25, 95% CI [0.94, 1.66],  $z = 1.54$ ,  $P = 0.124$ .

This suggests, as noted in the main article, that GPT-4 discovered more real articles than 3.5 *more generally*, and therefore was also more likely to discover relevant articles. Beyond this, however, it did not tailor article recommendations better to specific topics. However, when we exploratorily

examine this by topic breadth, the data tell a more nuanced story. By design, the Research Librarian task was somewhat easy at the highest broadness level: There were many influential papers, likely familiar to GPT, that met our criteria. On the other extreme, the narrowest subject areas were intentionally difficult. For example, few articles use the Implicit Association Test to predict dental outcomes, and they may justifiably be absent in GPT’s training data. The areas between these extremes – neither designed to be too easy nor too difficult – were thus of greater interest.

When limiting the analysis to non-hallucinations, GPT-4 shows relative improvement in finding relevant articles on Somewhat Broad and Moderate topics, with GPT-3.5’s success at generating relevant references dropping off more sharply as topics become narrower. As shown in Table S5, and reported in the main article, the difference between GPT-4 and GPT-3.5 on this count is significant for topics of moderate broadness, but not for topics at other broadness levels once correcting for multiple comparisons. (Since we are running 5 exploratory analyses here, a Bonferroni correction dictates that our significance level should be  $P < 0.01$ .)

**Table S5.** *Relevance Rates for Non-Hallucinations, by Broadness level*

| <b>Broadness</b> | <b>GPT-3.5 Proportion of Non-Hallucinations marked Relevant</b> | <b>GPT-4 Proportion of Non-Hallucinations marked Relevant</b> | <b>Odds Ratio</b> | <b>P-value</b> |
|------------------|-----------------------------------------------------------------|---------------------------------------------------------------|-------------------|----------------|
| Broad            | 75/80 (93.8%)                                                   | 99/100 (99.0%)                                                | 6.60              | 0.088          |
| Somewhat Broad   | 44/83 (53.0%)                                                   | 55/83 (66.3%)                                                 | 1.74              | 0.083          |
| Moderate         | 18/56 (32.1%)                                                   | 69/92 (75.0%)                                                 | 6.33              | < 0.001        |
| Somewhat Narrow  | 14/36 (38.9%)                                                   | 31/98 (31.6%)                                                 | 0.73              | 0.431          |
| Narrow           | 3/65 (4.62%)                                                    | 0/100 (0.0%)                                                  | 0                 | 0.03           |

Note: The  $P$ -value for Narrow topics is drawn from linear regression since GPT-4’s 0% disallows logistic regression. All other  $P$ -values are drawn from logistic regressions.

We also use logistic regression analysis to examine the significance of the difference in the curvilinear effect by GPT-Version. Because this analysis was exploratory, we ran the relevant regression analysis several ways as a robustness check. The first analysis begins by excluding hallucinations, since (as described above) not doing so tends to overestimate GPT-4's improvement in generating relevant articles. It captures a 3-way interaction, using logistic regression with Relevant (vs. Not Relevant) as the dependent variable, and 5 independent variables: GPT Version (4 vs. 3.5), Broadness, Broadness-square, GPT Version x Broadness, GPT Version x Broadness-square. The independent variables of particular interest here are GPT Version x Broadness and GPT Version x Broadness-square, since their significance indicates a different curvilinear relationship dependent on the version of GPT, in this case the fact that GPT-3.5's success (relative to GPT-4's) at locating relevant articles drops off more quickly as categories narrow. These coefficients are robustly significant. Specifically, the model is as follows: GPT Version ( $\beta = -3.590$ ,  $SE = 1.197$ ,  $z = -3.00$ ,  $P = 0.003$ ), Broadness ( $\beta = -0.010$ ,  $SE = 0.556$ ,  $z = -0.02$ ,  $P = 0.985$ ), Broadness-square ( $\beta = 0.177$ ,  $SE = 0.091$ ,  $z = 1.95$ ,  $P = 0.051$ ), GPT Version x Broadness ( $\beta = 2.947$ ,  $SE = 0.824$ ,  $z = 3.58$ ,  $P < 0.001$ ), and GPT Version x Broadness-square ( $\beta = -0.430$ ,  $SE = 0.131$ ,  $z = -3.27$ ,  $P = 0.001$ );  $Pseudo-R^2 = 0.3631$ . Next, instead of removing hallucinations we statistically controlled for them by adding them as a sixth independent variable in the regression. Though logistic regression is statistically preferable given the dichotomous variable, the perfect alignment between hallucination and irrelevance disallows this analysis, and so this was run as a regular linear regression. The critical interaction terms once again achieved robust significance in this regression. Specifically, the model is as

follows: GPT Version ( $\beta = -0.585$ ,  $SE = 0.108$ ,  $t(993) = -5.41$ ,  $P < 0.001$ ), Broadness ( $\beta = -0.015$ ,  $SE = 0.058$ ,  $t(993) = -0.26$ ,  $P = 0.794$ ), Broadness-square ( $\beta = 0.026$ ,  $SE = 0.010$ ,  $t(993) = 2.72$ ,  $P = 0.007$ ), GPT Version x Broadness ( $\beta = 0.420$ ,  $SE = 0.082$ ,  $t(993) = 5.15$ ,  $P < 0.001$ ), GPT Version x Broadness-square ( $\beta = -0.055$ ,  $SE = 0.013$ ,  $t(993) = -4.15$ ,  $P < 0.001$ ), and Hallucination vs. not ( $\beta = -0.435$ ,  $SE = 0.031$ ,  $t(993) = -14.12$ ,  $P < 0.001$ );  $R^2 = 0.4871$ . Third, we ran the initial logistic regression but added Incomplete (vs. complete) as a further independent variable. Again, the relevant interaction terms were robustly significant. Specifically, the model is as follows: GPT Version ( $\beta = -3.446$ ,  $SE = 1.214$ ,  $z = -2.84$ ,  $P = 0.005$ ), Broadness ( $\beta = 0.013$ ,  $SE = 0.558$ ,  $z = 0.02$ ,  $P = 0.982$ ), Broadness-square ( $\beta = 0.175$ ,  $SE = 0.091$ ,  $z = 1.92$ ,  $P = 0.055$ ), GPT Version x Broadness ( $\beta = 2.854$ ,  $SE = 0.834$ ,  $z = 3.42$ ,  $P = 0.001$ ), GPT Version x Broadness-square ( $\beta = -0.417$ ,  $SE = 0.133$ ,  $z = -3.15$ ,  $P = 0.002$ ), and Incomplete ( $\beta = 0.256$ ,  $SE = 0.379$ ,  $z = 0.67$ ,  $P = 0.500$ );  $Pseudo-R^2 = 0.3635$ . Fourth, we ran a second linear regression with the same additional IV (Incomplete vs. Complete). Again, the relevant terms were highly significant. Specifically, the model was: GPT Version ( $\beta = -0.574$ ,  $SE = 0.109$ ,  $t(992) = -5.27$ ,  $P < 0.001$ ), Broadness ( $\beta = -0.015$ ,  $SE = 0.058$ ,  $t(992) = -0.26$ ,  $P = 0.794$ ), Broadness-square ( $\beta = 0.026$ ,  $SE = 0.010$ ,  $t(992) = 2.74$ ,  $P = 0.006$ ), GPT Version x Broadness ( $\beta = 0.414$ ,  $SE = 0.082$ ,  $t(992) = 5.04$ ,  $P < 0.001$ ), GPT Version x Broadness-square ( $\beta = -0.055$ ,  $SE = 0.013$ ,  $t(992) = -4.07$ ,  $P < 0.001$ ), Hallucination vs. not ( $\beta = -0.444$ ,  $SE = 0.033$ ,  $t(992) = -13.49$ ,  $P < 0.001$ ), and Incomplete ( $\beta = 0.028$ ,  $SE = 0.037$ ,  $t(992) = 0.75$ ,  $P = 0.456$ );  $R^2 = 0.4874$ . We thus find that independent of modeling choices, the curvilinear interaction is highly significant: GPT-3.5 drops off more quickly than GPT-4 in terms of finding relevant articles as the topics narrow.

Relevant Citation Counts. In the main article, we reported that focusing only on relevant articles, GPT-4's selections did not have higher citation counts than those of GPT-3.5. Here, we begin by reporting the overall effects in greater detail. As noted in the main article, the two GPT models performed similarly here, with GPT-4's articles averaging 2,936.7 (95% CI [2128.3, 3745.1]) citations compared to 3,105.2 (95% CI [2205.5, 4005.0]) for GPT-3.5;  $t(406) = 0.2653$ ,  $P = 0.7909$ ,  $d = 0.027$ . It is possible, however, that a significant effect on this count was disguised by other factors. We reported, for example, that GPT-4 generated more relevant citations for somewhat broad and moderate topics. By pulling relevant articles in narrower domains, GPT-4 might be limited to less influential articles. In short, once we controlled for topic broadness, GPT-4 might have selected articles with higher citation counts. However, a regression analysis ruled out this possibility. Specifically, we limited our analysis to articles marked relevant and ran a simple linear regression with citation count as the dependent variable and GPT-4 (vs. GPT-3.5) and Topic Broadness as independent variables. While topic broadness was associated with higher citation count, the coefficient for GPT-version remained statistically insignificant in this model. Specifically, the model was: GPT Version ( $\beta = 218.38$ ,  $SE = 618.63$ ,  $t(405) = 0.35$ ,  $P = 0.724$ ), Broadness ( $\beta = 1512.33$ ,  $SE = 280.10$ ,  $t(405) = 5.40$ ,  $P < 0.001$ );  $R^2 = 0.0673$ .

Relatedly, GPT-4 might have proved superior at finding references that were important but also more recent. Such references would however have had less time to gather citations. To test for this possibility, we coded the year for the articles. We first ran a simple regression with article Year as the dependent variable and GPT-Version (4 vs. 3.5) as an independent variable, again limiting the analysis to only articles coded as "Relevant." This yielded a

modest but significant effect, the model being: GPT Version ( $\beta = 1.850$ ,  $SE = 0.821$ ,  $t(406) = 2.25$ ,  $P = 0.025$ );  $R^2 = 0.0124$ . This result suggests that GPT-4 pulled more recent articles than GPT-3.5. However, this result should be treated with caution for several reasons. First, the analysis was exploratory. Second, the effect was small. Third and most importantly, the effect disappears once we control for topic broadness. Adding this as an additional IV yields the model: GPT Version ( $\beta = 1.059$ ,  $SE = 0.755$ ,  $t(405) = 1.40$ ,  $P = 0.161$ ), Broadness ( $\beta = -3.091$ ,  $SE = 0.342$ ,  $t(405) = -9.05$ ,  $P < 0.001$ );  $R^2 = 0.1784$ . It appears that if GPT-4 indeed pulled slightly more current articles, this pattern was related to its tendency to better generate relevant references on narrower topics.

While not confident in the result that GPT-4 pulled more recent articles, we nevertheless ran a set of regressions to confirm that article aging was not disguising an effect of GPT-version on Citation Count. Toward this end, we ran two regression analyses. Again, limiting the analysis to articles coded as "Relevant," we first ran a regression with Citation Count as the dependent variable, and GPT-Version (4 versus 3.5) and article Year as independent variables. This again yielded an insignificant effect of GPT-Version on Citation Count, the model being: GPT Version ( $\beta = 737.23$ ,  $SE = 495.87$ ,  $t(405) = 1.49$ ,  $P = 0.138$ ), Article Year ( $\beta = -489.66$ ,  $SE = 29.80$ ,  $t(405) = -16.43$ ,  $P < 0.001$ );  $R^2 = 0.4001$ . Second, we ran a version of the regression that included both Topic Broadness and Year as covariates alongside GPT-Version. Once again, we see no significant partial effect of GPT-Version on article Citation Count. Specifically, the model was: GPT Version ( $\beta = 737.03$ ,  $SE = 497.96$ ,  $t(404) = 1.48$ ,  $P = 0.140$ ), Broadness ( $\beta = -1.26$ ,  $SE = 246.60$ ,  $t(404) = -0.01$ ,  $P = 0.996$ ), Article Year ( $\beta = -489.73$ ,  $SE = 32.71$ ,  $t(404) = -14.97$ ,  $P < 0.001$ );  $R^2 = 0.4001$ .

### Section S3, GPT as Research Ethicist – Detailed Methods

Full Research Design. In Study 2, we presented GPT-3.5 and GPT-4 with fictional vignettes describing flawed research protocols, posing as scientists looking for feedback on their work. These vignettes were very loosely based on real research projects that have proven difficult to replicate. For each, we described a series of plausible but non-obvious studies, but purposefully inserted evidence of 4 statistically unsound and ethically questionable practices. Three of these vignettes contained poor practices that were blatant. The researchers openly described a litany of scientific sins, such as examining data for statistical significance before deciding whether to continue collection, making post-hoc data exclusions that drew out significant effects, framing exploratory findings as if predicted in advance, and running many exploratory analyses but presenting only significant results without correcting for multiple comparisons.

Three further vignettes presented similar fictional protocols but indicated these problematic decisions more subtly. For example, the first matched pair of vignettes described a protocol to test for “emotional clairvoyance” (emotional contagion from interactions out of participants’ earshot). In the *blatant* version of this vignette, the researchers say, “after just 30 participants in each condition, effects already reached statistical significance ( $p < .05$ ), so we stopped data collection...” In the subtler version of this vignette, the researchers “collected 50 participants in each condition, at which point statistical analysis indicated that our results reached statistical significance ( $p < .05$ ).” They do not directly mention using this test to decide whether to collect more subjects, but later state that “In the second study, we were



worried we wouldn't always see such a large effect size, so we decided to collect data from 150 participants." Though there is nothing intrinsically wrong with collecting a larger sample in one study versus another, one familiar with problems of p-hacking might reasonably be suspicious. If they weren't stopping and starting collection, why was this decision made?

The purpose of this variation between blatant and subtle vignettes was to test, 1) whether GPT showed awareness of the clear methodological problems described in the blatant vignettes, and 2) whether it would be able to read between the lines and recognize violations in more realistically described research. In effect, the subtle vignettes were more ecologically valid descriptions, designed to better resemble how a real-world researcher might represent flawed research.

Each of our 6 vignettes was presented 36 times (18 times each to GPT-3.5 and GPT-4). As detailed later, we exploratorily varied the initial prompt we used to request ChatGPT's feedback on these protocols. In addition, we ran a secondary study (described later) where ChatGPT responded to versions of two of these vignettes that contrastingly described pristine ethics and statistical practices.

For each vignette in the main study, a rubric was created, containing 10 points we wanted to see GPT highlight in each chat. Four of the ten points were for recognizing problematic research practices. For each of these four, a second point could be received by suggesting a practice that remedied the problem. GPT might receive these "solution points" for a sufficiently strong statement of the problem. For example, GPT might generate a response like "You have started and stopped collection part way through to test for

significance. This is a form of p-hacking and should not be done since it inflates the likelihood of false positive results.” The solution in this case is intrinsic to the statement of the problem: Do not do this. This line would therefore receive points for both the statement of the problem and the solution. On the other hand, if GPT said something like “You should decide your sample size in advance based on a power analysis,” this would receive credit for proposing a solution, but in the absence of further statements would not receive credit for pointing out the problem. In addition to these 8 possible points, GPT was credited if it suggested 1. Pre-registration or other firm pre-research decisions, and 2. Replication.

Detail on vignettes and their coding. The full vignettes and the rubrics used to code them may be found in the document “GPT as Research Ethicist Materials 20240227” at <https://osf.io/sdahr/>. Each vignette was written by the authors and designed to contain evidence (blatant or subtle) of 4 substandard and ethically questionable research practices. The first matched pair of vignettes contained descriptions of research on “Emotional Clairvoyance,” in which participants supposedly experienced emotional contagion from interactions that took place out of earshot. The second set of vignettes described studies of “Atmospheric Cognition,” in which subtle characteristics of the space one is in impact mental states and behaviors. The third set of vignettes described an econometrical study, using data from a large dating site to demonstrate that more extroverted individuals are perceived as less attractive. Table S6 displays a list of the problematic practices found in these vignettes, and how they varied between the Blatant and Subtle Versions.

**Table S6.** *Details of ethical/methodological violations described in vignettes*

|                                                                            | <b>Blatant Version</b>                                                                                                                                                                                                                       | <b>Subtle Version</b>                                                                                                                                                                                                                                              |
|----------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Vignette 1a &amp; 1b: Emotional Clairvoyance</b>                        |                                                                                                                                                                                                                                              |                                                                                                                                                                                                                                                                    |
| Issue 1: Stopping & Starting Data Collection to Test Significance          | <ul style="list-style-type: none"> <li>• Tested for significance and then stopped collection after 30 subjects.</li> <li>• Later, collected additional subjects when significance was not initially reached.</li> </ul>                      | <ul style="list-style-type: none"> <li>• Collected notably different sample sizes in different studies.</li> <li>• In final study, indicated “we decided to collect additional data,” with ambiguity around whether significance testing played a role.</li> </ul> |
| Issue 2: Testing Several Moderators w/o Correction Multiple Comparisons    | <ul style="list-style-type: none"> <li>• Tested the Big 5 personality traits as moderators, without mentioning correction for multiple comparisons.</li> </ul>                                                                               | <ul style="list-style-type: none"> <li>• Tested the Big 5 personality traits as moderators, without mentioning correction for multiple comparisons.</li> </ul>                                                                                                     |
| Issue 3: Hypothesizing After Results are Known (HARKing)                   | <ul style="list-style-type: none"> <li>• Says moderator from exploratory analysis makes intuitive sense and will be reported as what they were testing at the outset.</li> </ul>                                                             | <ul style="list-style-type: none"> <li>• Says moderator from exploratory analysis makes intuitive sense but does not indicate whether or not it will be framed as exploratory.</li> </ul>                                                                          |
| Issue 4: Selective Reporting                                               | <ul style="list-style-type: none"> <li>• Attains null results in two studies with slightly different conditions. Indicates they will not publish them.</li> </ul>                                                                            | <ul style="list-style-type: none"> <li>• Attains null results in one study with slightly different conditions. Indicates uncertainty about the study, and that they may “focus” on the other conditions.</li> </ul>                                                |
| <b>Vignette 2a &amp; 2b: Atmospheric Cognition</b>                         |                                                                                                                                                                                                                                              |                                                                                                                                                                                                                                                                    |
| Issue 1: Suppressing Unsuccessful Pilot Studies                            | <ul style="list-style-type: none"> <li>• Says they piloted “a bunch of versions” with initially mixed results.</li> <li>• In the main version, piloted a couple of versions that were too subtle, which they therefore “ignored.”</li> </ul> | <ul style="list-style-type: none"> <li>• Says they initially piloted a “few versions” with initially mixed results. Whether they will share these mixed results is left ambiguous.</li> </ul>                                                                      |
| Issue 2: Post-hoc Exclusion of Participants to Achieve Significance        | <ul style="list-style-type: none"> <li>• Said they removed math majors who “skewed the results,” only after which results achieved significance.</li> </ul>                                                                                  | <ul style="list-style-type: none"> <li>• Said they noticed and removed math majors who had an advantage and were “unevenly assigned” and that results were significant.</li> </ul>                                                                                 |
| Issue 3: Adding Further Subjects When Significant Results are Not Attained | <ul style="list-style-type: none"> <li>• Openly said that when they didn't see an effect after 100 participants in Study 3, they added 50 more to reach significance.</li> </ul>                                                             | <ul style="list-style-type: none"> <li>• Said they realized they would need more subjects in Study 3, and so collected 50 more. Does not directly indicate they tested significance.</li> </ul>                                                                    |
| Issue 4: Reporting Marginal Results as if They Reached Significance        | <ul style="list-style-type: none"> <li>• Said they achieved a significance level of <math>p=.052</math>, “which rounds down to .05”.</li> </ul>                                                                                              | <ul style="list-style-type: none"> <li>• Said they achieved a significance level of <math>p=.052</math>, “which is so close to the <math>p=.05</math> threshold as to be statistically significant for most...purposes.”</li> </ul>                                |

| <b>Vignette 3a &amp; 3b: Bad Econometrics</b>                                                   |                                                                                                                                                                                                      |                                                                                                                                                                                   |
|-------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Issue 1: Suppressing Outcome Variables That Yielded Null Results                                | <ul style="list-style-type: none"> <li>Because two alternative DVs didn't yield significant results, they "ignored these" and feel they can be "omitted from the report."</li> </ul>                 | <ul style="list-style-type: none"> <li>Since two alternative DVs didn't yield significant results, they were "downplayed," and the report will "focus" on the third.</li> </ul>   |
| Issue 2: Excluding Parts of Data that Did Not Yield Significant Results                         | <ul style="list-style-type: none"> <li>Exploratorily limited analyses to people under 60 and who live in specific regions, because this yielded stronger results.</li> </ul>                         | <ul style="list-style-type: none"> <li>Exploratorily limited analyses to people under 60, the primary demographic of the website, which made results stronger.</li> </ul>         |
| Issue 3: Cycling Through Covariates to Attain Results, w/o Corrections for Multiple Comparisons | <ul style="list-style-type: none"> <li>Indicated that they had the "Big 5 and dozens of subtler measures", and later "examined various potential covariates" to find significant model.</li> </ul>   | <ul style="list-style-type: none"> <li>Indicated that they had the "Big 5 and subtler measures", and later that they "examined potential covariates."</li> </ul>                  |
| Issue 4: Violating Assumptions Around Multicollinearity & Overfitting                           | <ul style="list-style-type: none"> <li>Includes two variables with a high likelihood of multicollinearity (Shyness, Self-Confidence) in model to show significant effect of Extroversion.</li> </ul> | <ul style="list-style-type: none"> <li>Includes one variable with potential for multicollinearity (Assertiveness) in model to show significant effect of Extroversion.</li> </ul> |

The same rubrics were used for scoring the Blatant and Subtle versions of these vignettes. These rubrics were as follows, with GPT given a point by the coders for each affirmative answer.

Coding Rubrics for Scientific Ethicist Study. Below are the full rubrics used to code responses to the Scientific Ethicist vignettes.

Rubric for Vignettes 1a & 1b:

1. Did it draw attention to the issue of making decisions about data collection (stopping/starting) by examining data part-way through collection?
2. Did it provide advice (e.g. power analysis, sticking to a pre-decided sample size, not testing part-way through collection) that mitigates the issues around stopping/starting data collection based on data examination during collection?

3. Did it draw attention to the issue of testing different potential variables/moderators, creating problems of multiple comparisons or increasing potential for type 1 errors?
4. Did it provide advice (e.g. corrections for multiple comparisons, deciding hypothesized moderators in advance) that mitigates the issues around multiple comparisons?
5. Did it draw attention to the issue of HARKing or otherwise ignoring the exploratory nature of certain findings?
6. Did it provide advice (e.g. clear delineation of exploratory analyses, pre-registering hypotheses) that mitigates the problem of HARKing?
7. Did it draw attention to the problem of suppressing null results in reporting?
8. Did it provide advice (e.g. reporting all results) that mitigates the issue of suppressing null results?
9. Did it suggest further replication of results?
10. Did it suggest pre-registration or other forms of commitment to pre-decided methodological strategies?

Rubric for Vignettes 2a & 2b:

1. Did it draw attention to the issue of ignoring or suppressing pilot studies?
2. Did it provide advice (e.g. discussing pilot research, publishing all results) that mitigates the issue of ignoring or suppressing pilot studies?
3. Did it draw attention to the issue of making inappropriate post hoc exclusions of participants?
4. Did it provide advice (e.g. not making exclusions, deciding exclusions in advance) that mitigates the issue of inappropriate post hoc exclusions?
5. Did it draw attention to the issue of adding further subjects to studies when statistically significant results are not initially attained?
6. Did it provide advice (e.g. power analysis, sticking to pre-decided sample, not testing part-way through collection) that mitigates the issue of adding subjects to a study when not seeing significant results?
7. Did it draw attention to the issue of reporting marginal results as if they'd met thresholds for significance?

8. Did it provide advice (e.g. reporting exact p-values, not rounding down) that mitigates the issue or reporting marginal results as if they were significant?
9. Did it suggest further replication of results?
10. Did it suggest pre-registration or other forms of commitment to pre-decided methodological strategies?

Rubric for Vignettes 3a & 3b:

1. Did it draw attention to the issue of ignoring or suppressing outcome variables that didn't lead to favorable results?
2. Did it provide advice (e.g. reporting on all possible outcome variables) that mitigates the issue of ignoring or suppressing unfavorable outcome variables?
3. Did it draw attention to the issue of excluding chunks of data that did not fit results?
4. Did it provide advice (e.g. including all data, deciding exclusions a-priori and/or clearly reporting and defending exploratory exclusions) that mitigates the issue of inappropriately excluding data?
5. Did it draw attention to the issue of cycling through many potential covariates until finding ones that yielded desirable results?
6. Did it provide advice (e.g. corrections for multiple comparisons, hypothesizing covariates in advance) that mitigates the issue of examining many potential covariates?
7. Did it draw attention to the issue of potentially violating statistical assumptions (e.g. issues around overfitting, multicollinearity)?
8. Did it provide advice (e.g. split samples, checks for multicollinearity) that mitigates issues around potentially violating statistical assumptions?
9. Did it suggest replications of results in other data or contexts and/or urge caution about result generalizability?
10. Did it suggest pre-registration or other forms of commitment to pre-decided methodological strategies?

Coding of Vignettes. Responses to the vignettes were blinded by prompt condition. Where possible, content that would make the initial prompt obvious (but was otherwise irrelevant) was removed. Coders were also

blinded to whether responses were given by GPT-3.5 or GPT-4, as well as to whether each chat was in response to a blatant or subtle version of the vignette. GPT's responses were placed in groups coinciding with the topic of the vignette but were otherwise presented in a randomized order.

Two coders independently rated each GPT response on each of the 10 points from the relevant rubric. Points of disagreement were discussed. Following discussion, responses were averaged, with GPT receiving partial credit in cases where one but not the other coder felt a point was merited. Inter-rater coding achieved a high degree of reliability, with a Cronbach alpha of  $\alpha = 0.9827$ . Of 2,160 possible points GPT could receive, the two coders disagreed on 34 (~1.57%).

One coder was an author, and the other was a Research Assistant who was kept blind to the full purpose of the research until after completion. Responses were first coded independently by each coder. Points of disagreement were then discussed, with the coder who was not an author encouraged to lead discussions to minimize pressure to conform to the author's decision. Though the rubrics were decided in advance, certain points needed to be clarified early in the coding. Specifically, after discussing a few instances of question 10 (pre-registration or other firm methodological commitments), the coders made the decision that credit would be given for either recommending preregistration or making *at least two* other recommendations around committing to decisions in advance. This was to prevent GPT from too easily receiving credit for the same point multiple times. For example, recommending a power analysis to decide on sample size in advance received credit for q2 in Vignettes 1a-b. To also receive credit

for q10, GPT would need to make a second point around deciding methodologies in advance. For q7 in Vignettes 3a-b, two decisions were made. First GPT was credited for suggesting that the vignette's interpretation of their effects might be flawed due to complicated relationships (e.g. mediating roles) between the two variables, since this point was conceptually very similar to the one we wished it to make: The fundamental problem was the complex and likely strong relationships between the different independent variables. On the other hand, GPT was not credited on this question for answers highlighting statistical assumptions that were not clearly violated (e.g. homoskedasticity, linearity), as such observations seemed generic and did not acknowledge the fundamental problem (overfitting, multicollinearity) in the regression analysis.

Variation in Initial Prompts. As discussed in the main text, we varied the initial prompt used when initiating each chat with GPT, to test whether this would impact GPT's performance on reviewing the vignettes. In this section, we provide more detail around these initial prompts. Complete prompt stimuli may be found in the document "GPT as Research Ethicist Materials 20240227" at <https://osf.io/sdahr/>.

1. Level of Feedback, Three basic requests varied in the level of requested feedback, from simply asking for "feedback" to requesting GPT "give a detailed response and be as critical as possible about any flaws you notice." The goal was to ascertain GPT might give a more rigorous response when asked directly for more critical feedback.
2. Theory Protectiveness. In these prompts, we indicated that we either are or are not "deeply invested" in our theory. The goal was to ascertain whether GPT would be "kinder" and thus less critical when we indicated deep commitment to the theory.



3. **Researcher Status:** In one version, we pretended to be a novice researcher running our first project. In the other, we claimed to be an award-winning researcher with a tenured chair at an elite university. The goal was to ascertain whether GPT would provide less critical feedback to a researcher who was presented as higher status and thus potentially more knowledgeable.
4. **Pro- and Anti-Open Science.** In one version, we claimed to believe deeply in open science practices. In the second we prefaced our request with an argument that the movement against “p-hacking” is overblown and slows down the march of science. The goal was to ascertain whether GPT would give stronger feedback following a pro-open science argument.
5. **Conciseness.** In one version, we asked for a very concise response. In the other, we asked GPT to disregard length limitations and write as long a response as possible. The goal was to ascertain whether GPT would be less successful in its feedback when asked to be more concise.
6. **Rejection Sensitivity.** In one version, we claimed to be very sensitive to criticism. In the second, we asked GPT not to worry about being kind in its responses. The goal was to ascertain whether GPT would be gentler (and thus less critical) when the person to whom it is speaking appears to have difficulty taking criticism.
7. **Jailbreaking version 1.** In one version, we asked GPT to “carefully consider issues around p-hacking and open science.” In a second version (7a), we asked it to “pretend to be a person who is not at all concerned with issues like p-hacking and open science.” In a third (7b), we asked it to “not include any mentions of p-hacking in your feedback.” The goal was to ascertain whether GPT would give less ethical feedback when directly instructed to do so.
8. **Jailbreaking version 2.** In one version, we asked GPT to impersonate Uri Simonsohn, a well-known advocate of open science and replicable statistical practices. In a second, we asked it to impersonate Diederik Stapel, an infamous p-hacker and data fabricator, who has, at the time of writing, had 58 articles retracted according to [retractionwatch.com](https://retractionwatch.com). The goal was to ascertain whether GPT would give less ethical feedback when asked to impersonate an individual known for poor ethical conduct compared to one known for strong ethical conduct.

Each of these effectively contained a “Good Performance Prompt”, hypothesized to elicit a more rigorous response from GPT, and a “Bad Performance Prompt”, designed to elicit a less rigorous response. For example, it was possible that GPT would be more likely to call out p-hacking when our initial request was prefaced by a pro- (relative to an anti-) open science argument. In separate chats, we presented GPT-3.5 and GPT-4 with each of these 18 initial prompts, followed by each of our 6 vignettes. We thus received and coded a total of 216 responses, each of which was judged on our 10-point rubric. Full transcripts of GPT’s responses to our vignettes may be found in the document “GPT as Research Ethicist Transcripts 20240220” at <https://osf.io/sdahr/>.

#### **Section S4: GPT as Research Ethicist – Supporting Analyses**

In this section, we will present more detailed and alternative analyses of the Research Ethicist study. Except when otherwise noted (e.g. the non-parametric tests), contrasts were conducted using t-tests. All analyses were conducted in Stata.

More Detailed Overall Results. In the Blatant condition, GPT-4 achieved a mean score of 8.86 out of a possible 10 points (95% CI 8.55, 9.18), while GPT-3.5 averaged 5.39 (95% CI [4.79, 5.98]). This difference was significant and large,  $t(106) = -10.3511$ ,  $P < 0.0001$ ,  $d = 1.992$ . Similarly, in the Subtle condition, GPT-4 averaged 7.26 points (95% CI [6.72, 7.80]) compared to GPT-3.5’s 4.05 (95% CI [3.47, 4.62]);  $t(106) = -8.1619$ ,  $P < 0.0001$ ,  $d = 1.571$ .

GPT-4 scored higher on blatant relative to subtle vignettes ( $t(106) = 5.1297$ ,  $P < 0.0001$ ,  $d = 0.987$ ), as did GPT-3.5 ( $t(106) = 3.2565$ ,  $P = 0.0015$ ,  $d = 0.627$ ).

Finally, GPT-4 received more points in response to subtle vignettes than GPT-3.5 did in response to blatant ones ( $t(106) = -4.6618$ ,  $P < 0.0001$ ,  $d = 0.897$ ).

Detailed Analysis of Initial Prompts. The variation of initial prompts was exploratory. Furthermore, since each separate prompt was used only 12 times, and we analyzed data on the level of the chat, statistical power for individual prompts was low. This analysis was therefore designed to pick up only large effects: Does GPT do much better after certain prompts than after others? We completed 12 exploratory analyses. We conservatively used a Bonferroni correction for multiple comparisons, and so our threshold for statistical significance was  $P < 0.004167$ .

We began by examining each matched pair of prompts (e.g. Low versus High Researcher Status, or Pro- versus Anti-Open Science). As seen in Table S7, most of these basic comparisons did not approach statistical significance. As noted also in the main article, one comparison was a tentative exception: Requesting GPT to “carefully consider issues around p-hacking and open science” ( $M = 8.38$ , 95% CI [7.64, 9.11]) elicited a stronger response than requesting that it “not include any mentions of p-hacking” ( $M = 6.08$ , 95% CI [4.43, 7.74];  $t(22) = -2.7912$ ,  $P = 0.0106$ ,  $d = 1.140$ ). However, this result should be treated with caution considering 1) the small sample size, and 2) that it did not meet the more stringent significance threshold dictated by the Bonferroni correction. No paired comparison reached statistical significance after adjusting for multiple comparisons.

**Table S7.** *Response Quality by Initial Prompt*

| <b>Comparison</b>                                            | <b>Mean (out of 10):<br/>Following Good<br/>Performance<br/>Prompt</b> | <b>Mean (out of 10):<br/>Following Bad<br/>Performance<br/>Prompt</b> | <b>Statistical<br/>Significance of<br/>Difference</b> |
|--------------------------------------------------------------|------------------------------------------------------------------------|-----------------------------------------------------------------------|-------------------------------------------------------|
| <i>Variation of Level of Feedback Requested</i>              |                                                                        |                                                                       |                                                       |
| High vs. Low/Basic Feedback Request                          | 6.42                                                                   | 5.90                                                                  | $P = 0.6055$                                          |
| Low vs. High Conciseness                                     | 5.63                                                                   | 6.29                                                                  | $P = 0.5609$                                          |
| <i>Manipulations of Perceived Researcher Characteristics</i> |                                                                        |                                                                       |                                                       |
| Low vs. High Theory Protectiveness                           | 5.54                                                                   | 5.50                                                                  | $P = 0.9744$                                          |
| Low vs. High Researcher Status                               | 5.33                                                                   | 5.75                                                                  | $P = 0.7285$                                          |
| Low vs. High Rejection Sensitivity                           | 5.75                                                                   | 5.54                                                                  | $P = 0.8347$                                          |
| <i>Encouragement of More or Less Ethical Responses</i>       |                                                                        |                                                                       |                                                       |
| Pro- vs. Anti-Open Science                                   | 7.50                                                                   | 7.08                                                                  | $P = 0.6187$                                          |
| Concerned w/ p-hacking vs. Not Concerned                     | 8.38                                                                   | 7.08                                                                  | $P = 0.0845$                                          |
| Concerned w/ p-hacking vs. Don't Mention                     | 8.38                                                                   | 6.08                                                                  | $P = 0.0106$                                          |
| Ethicist vs. Data Fabricator                                 | 7.88                                                                   | 7.46                                                                  | $P = 0.6470$                                          |

Notes: The first category listed in the Comparison column was hypothesized as the “Good Performance Prompt” in each pair.

Though individual pairs of prompts did not elicit statistically distinct responses, we conducted three more analyses that aggregated groups of responses in meaningful ways. First, we asked whether the three “jailbreaking” prompts, in which we specifically asked GPT to act in an unethical manner (Pretend not to be concerned with p-hacking; Do not mention p-hacking; Impersonate a known data fabricator) elicited weaker responses than their “Anti-jailbreak” control prompts (Carefully consider issues around p-hacking and open science; Impersonate an ethicist). While the two “Anti-jailbreaking” responses elicited descriptively stronger responses ( $M = 8.13$ , 95% CI [7.43, 8.82]) than the three “jailbreaking” responses ( $M = 6.88$ , 95% CI [6.07, 7.68]), the difference did not reach

significance after controlling for multiple comparisons;  $t(58) = -2.2326$ ,  $P = 0.0295$ ,  $d = 0.588$ .

Second, we noticed that regardless of how we asked GPT to behave, responses looked descriptively stronger across prompts where we in any way evoked data ethics. For example, notice in Table S7 that the 12 chats where we asked GPT to impersonate a known data fabricator – designed as a jailbreaking prompt to elicit an unethical response – actually yielded responses that were descriptively higher-quality than most. Accordingly, we collapsed responses across all prompts that in any fashion evoked data ethics (Pro- vs. Anti-Open Science; Concerned vs. Not Concerned with p-hacking; Don't mention p-hacking; data ethicist vs. data fabricator), comparing these to all the remaining prompts, which contained no mentions of p-hacking or open science. Here, we found a statistically robust difference. GPT provided higher-quality responses after prompts that evoked data ethics ( $M = 7.35$ , 95% CI [6.89, 7.81]) than after prompts that did not ( $M = 5.78$ , 95% CI [5.30, 6.25]);  $t(214) = -4.4761$ ,  $P < 0.0001$ ;  $d = 0.625$ .

Considering the last result, we decided to conduct one more exploratory comparison as a robustness check. We asked whether the effect of evoking data ethics was sufficiently powerful that it would survive even when we limited initial prompts to those in which the context involved us specifically encouraging GPT to behave badly. In other words, would GPT perform better following the 4 prompts where we cast data ethics in a negative light (Anti-open Science argument; Pretend not to be concerned with p-hacking; Don't mention p-hacking; Impersonate known Data Fabricator) compared to those that did not evoke data ethics at all? Indeed, GPT's responses were higher-

quality following prompts that cast a negative light on research ethics ( $M = 6.93$ , 95% CI [6.28, 7.58]) compared to those that did not evoke research ethics ( $M = 5.78$ , 95% CI [5.30, 6.25]);  $t(178) = -2.6078$ ,  $P = 0.0099$ ,  $d = 0.440$ . While this result did not survive a Bonferroni correction, we argue that it is nevertheless meaningful. The effect is in the *opposite* direction from one we might have predicted a-priori, and this comparison was selected specifically as a tougher test of the effect of priming data ethics. At the least, it seems likely that the positive effect of evoking data ethics is a real one, since we see a strongly trending effect even when this prime is placed within a context of attempting to jailbreak GPT to elicit a poor response.

Non-Parametric Tests. More typical analyses (using means, t-tests, Cohen's d, etc.) may be easier for many researchers to understand and interpret.

However, the data from GPT-4 in Study 2 violated assumptions of normal distribution, with most responses skewed towards the higher end of our rubrics. To account for this and corroborate our findings, we report non-parametric analyses (Wilcoxon Rank-Sum Tests) and median averages, as an additional statistical treatment of these data. The effect sizes reported below are rank-biserial correlations based on these tests. The Wilcoxon Rank-Sum tests were run in Stata, and the correlations based on them were calculated in Excel.

We first examine the relative success of GPT-4 vs. GPT-3.5 in identifying and correcting ethical/statistical violations. For the Blatant vignettes, a Wilcoxon Rank-Sum test revealed statistically significant differences between GPT-4 (Median = 9) and GPT-3.5 (Median = 5),  $z = 7.363$ ,  $P < 0.001$ , with a large effect size ( $r = 0.709$ ). For Subtle Vignettes, as well, GPT-4 (Median = 8) outscored

GPT-3.5 (Median = 4);  $z = 6.522$ ,  $P < 0.001$ ,  $r = 0.628$ . GPT-4 scored higher in responses to Blatant (Median = 9) vs. Subtle (Median = 8) vignettes;  $z = 4.949$ ,  $P < 0.001$ ,  $r = 0.476$ . GPT-3.5 also scored higher in response to Blatant (Median = 5) vs. Subtle (Median = 4) vignettes;  $z = 3.152$ ,  $P = 0.002$ ,  $r = 0.303$ . Finally, GPT-4 scored higher on Subtle vignettes (Median = 8) compared to GPT-3.5 on Blatant vignettes (Median = 5), indicating the high degree of GPT-4's superiority on this task;  $z = 4.484$ ,  $P < 0.001$ ,  $r = 0.431$ .

We also conducted Wilcoxon Rank-Sum tests to explore the effects of different initial Positive vs. Negative performance prompts, as shown in Table S8. As with the analyses reported in the main text, these were exploratory and thus should undergo corrections for multiple comparisons. A conservative Bonferroni correction dictates that our significance threshold should be  $P < 0.004167$ . Therefore, none of the 9 contrasts in Table S8 reaches significance after this correction. The one strong trend (better performance when asked to review the vignettes as one concerned about p-hacking, versus asked not to mention p-hacking in the response) may indicate a jailbreaking effect but should be interpreted with significant caution considering the small sample size and lack of significance after correction for multiple comparisons.

**Table S8.** *Effects of Initial Prompts, Wilcoxon Rank-Sum tests*

| Comparison                                  | Median (out of 10):<br>Following Good<br>Performance<br>Prompt | Median (out of 10):<br>Following Bad<br>Performance<br>Prompt | Correlation  | Statistical<br>Significance |
|---------------------------------------------|----------------------------------------------------------------|---------------------------------------------------------------|--------------|-----------------------------|
| High vs. Low/Basic<br>Feedback Request      | 6                                                              | 5.5                                                           | $r = 0.062$  | $P = 0.711$                 |
| Low vs. High Conciseness                    | 6.5                                                            | 6.25                                                          | $r = -0.136$ | $P = 0.505$                 |
| Low vs. High Theory<br>Protectiveness       | 6                                                              | 5                                                             | $r = 0.018$  | $P = 0.931$                 |
| Low vs. High Researcher<br>Status           | 4                                                              | 6                                                             | $r = -0.024$ | $P = 0.907$                 |
| Low vs. High Rejection<br>Sensitivity       | 6                                                              | 5.5                                                           | $r = 0.030$  | $P = 0.884$                 |
| Pro- vs. Anti-Open Science                  | 7.5                                                            | 7.25                                                          | $r = 0.108$  | $P = 0.598$                 |
| Concerned w/ p-hacking<br>vs. Not Concerned | 8.75                                                           | 7.5                                                           | $r = 0.289$  | $P = 0.157$                 |
| Concerned w/ p-hacking<br>vs. Don't Mention | 8.75                                                           | 7                                                             | $r = 0.467$  | $P = 0.022$                 |
| Simonsohn vs. Stapel                        | 9                                                              | 8                                                             | $r = 0.102$  | $P = 0.619$                 |

Notes: The first category listed in the Comparison column was hypothesized as the “Good Performance Prompt” in each pair.

Finally, aligning with the main text of this article, we conducted three additional exploratory analyses. First, we combined data from the two “Anti-Jailbreaking” prompts to see if the responses that followed were superior to those following the three “Jailbreaking” prompts. A Wilcoxon Rank-Sum Test reveals that responses scored descriptively higher on our rubric after the Anti-Jailbreaking prompts (Median = 9) than after the Jailbreaking prompts (Median = 8), however the contrast did not reach statistical significance after correcting for multiple comparisons;  $z = 2.037$ ,  $P = 0.042$ ,  $r = 0.263$ .



Next, we examined whether there was an overall effect of priming data ethics, contrasting prompts that in any way evoked these (pro- and anti-open science, concerned w/ or not concerned with p-hacking, don't mention p-hacking, impersonate ethical or unethical researcher) compared to all other initial prompts. This contrast was statistically significant, even after correcting for multiple comparisons. Specifically, responses were superior following those prompts that in any way evoke data ethics (Median = 8) compared to those that did not (Median = 6);  $z = 4.134$ ,  $P < 0.001$ ,  $r = 0.281$ . Finally, as a robustness check, we conducted the same analysis but limited the prompts that evoked data ethics only to those that did so in a *negative* fashion, e.g. requesting that GPT generate text like somebody who is not concerned about p-hacking or impersonate a real-world unethical researcher. Even when limiting our analysis to prompts that were designed to deter GPT from giving a good response, this difference appears, though (as before) it does not survive the Bonferroni correction. Specifically, GPT gave superior responses following prompts that *negatively* evoked data ethics (Median = 8) compared to those that did not evoke data ethics at all (Median = 6);  $z = 2.482$ ,  $P = 0.0131$ ,  $r = 0.185$ .

### **Section S5: Good Research Vignette Study**

As described briefly in the main article, we conducted a second briefer study designed to examine the positive case in our Scientific Ethicist work. Specifically, we generated two new fictional vignettes, analogous to Vignettes 1a-b and Vignettes 2a-b in our main study. In contrast to these vignettes, however, Vignettes 1c and 2c describe research that is conducted with a very high degree of experimental and statistical rigor, and that otherwise displays

pristine practices around data ethics and open science. In these vignettes, all exploratory analyses are framed as such and undergo correction for multiple comparisons. Where possible, studies are pre-registered and replicated. All materials are made public. Moreover, these protocols describe practices that go above and beyond typical good research practices. For example, the researchers - recognizing the work will be contentious - do things like having rival researchers with competing hypotheses directly replicate their results or hiring statistical consultants to verify the properness of their analytical methodologies. GPT-4 and GPT-3.5 were each presented with each vignette 30 times for a total of 120 responses. The initial prompts were not varied in this study: In each conversation, ChatGPT was asked to identify the good research practices that stand out in the protocol. (Full vignettes and coding rubrics may be found in the document "GPT as Research Ethicist Materials 20240227" and full transcripts in the document "GPT as Research Ethicist Transcripts 20240220", both at <https://osf.io/sdahr/>.)

As in the main Scientific Ethicist Study, each response was coded based on a 10-point rubric. The rubrics described different practices that stood out as rigorous and ethical in the research protocols. Two coders examined the 120 transcripts, which were blinded to the version of GPT. As in the main study, each coder reviewed all 120 transcripts, giving 1 point for each practice from the rubric that GPT identified. Points of disagreement were then discussed, with the coder who was not an author on the paper asked to lead this discussion. In cases where the coders disagreed following this discussion, their two responses were composited. The final responses of the two coders achieved a very high degree of reliability; Cronbach's  $\alpha = 0.9695$ . In total, the coders disagreed on 11/1200 (~0.92%) of possible points ChatGPT could

receive. (Full data may be found in the spreadsheet “Good Research Ethicist Data Collapsed 20240227” at <https://osf.io/sdahr/>.)

For analysis, data were aggregated to the level of the chat, with each response receiving a minimum of 0 points and a maximum of 10. We report means for clarity of interpretation, but conduct (non-parametric) Wilcoxon Rank-Sum tests for gauging significance, since the results are not normally distributed.

Results. Both GPT-4 and GPT-3.5 did well on this task. In total, GPT-4 averaged 92.67% of possible points ( $M = 9.267$ ) compared to 90.42% ( $M = 9.042$ ) for GPT-3.5. A Wilcoxon Rank-Sum test reveals similar performance for GPT-4 (Median = 9) compared to GPT-3.5 (Median = 9) with the contrast trending toward but failing to reach statistical significance;  $z = 1.806$ ,  $P = 0.071$ ,  $r = 0.165$ . Exploratory analysis gauged the separate performance of GPT-4 vs. GPT-3.5 on each of the two vignettes. In response to Vignette 1c, GPT-4 ( $M = 9.533$ , Median = 10) slightly outperformed GPT-3.5 ( $M = 9.217$ , Median = 9), but the difference failed to meet the adjusted threshold of  $P < 0.025$  dictated by a Bonferroni correction;  $z = 2.114$ ,  $P = 0.035$ ,  $r = 0.273$ . For Vignette 2c, the performance of GPT-4 ( $M = 9.000$ , Median = 9) and GPT-3.5 ( $M = 8.867$ , Median = 9) were statistically indistinguishable;  $z = 0.622$ ,  $P = 0.534$ ,  $r = 0.080$ .

In contrast to the main study, where GPT-4 did sharply better than GPT-3.5 at identifying and correcting poor ethical and statistical practices, both proved highly competent at identifying *good* research practices. The reason for this distinction is somewhat puzzling. It is possible that the task was simply easier. In line with that interpretation, there was no “Subtle” condition in this

study: The vignettes are more analogous to the Blatant condition, on which GPT performed better. However, in the bad vignette study, GPT-3.5 failed to perform well even on the Blatant vignettes. An alternative interpretation is that GPT-3.5 is tuned to be somewhat “nicer” than GPT-4. Asked to review bad and good research protocols alike, it often tended to find good things to say, thereby succeeding in reviewing good protocols but failing at reviewing bad ones. GPT-4, in contrast, shows greater critical ability: In reviewing good research protocols, it was successful at drawing out and describing the good research practices. But when reviewing bad protocols, it was better able to be critical, nearly as easily recognizing and correcting *bad* practices. Finally, the result may be an artifact of the specific prompts used in this study. For example, GPT-4 might have demonstrated a higher rate of true positive detection had we used prompts that neutrally requested feedback, as opposed to explicitly requesting it identify good practices. Future research should tease apart these competing possibilities.

## **Section S6: GPT as Data Generator – Detailed Methods**

Full Research Design. As detailed in the main article, we used data from ChatGPT to attempt to replicate 4 common and well-studied stereotypes: Gender Attitudes, Gender Art/Science stereotypes, Gender Home/Work stereotypes, and Gender Math/Reading stereotypes. Adapting the stimuli used by Charlesworth and colleagues (1), we presented GPT-3.5 and GPT-4 with randomly ordered word dyads, requesting it estimate the cultural associations between each based on its training data. These dyads contained every possible combination of one word related to one of the target categories (the groups being stereotyped, in this case “Male” or “Female”)

with a second word related to one of the attribute categories (the characteristics being associated with these groups, e.g. “Home” or “Work”). We requested that, based on the cultural knowledge reflected in its training data, GPT estimate the overall cultural association between each pair of words. GPT was instructed to present each association on a scale of 1 to 10, with 1 meaning “the two concepts have a very low association with each other, based on the two words rarely being seen together” and 10 meaning “the two concepts have a very high association with each other, based on the words frequently being seen together.” To ensure responses were reasonably continuous, GPT was instructed (and, if needed, reminded) to round answers to the nearest tenth (e.g. 7.1) rather than to whole numbers. To preserve a consistent scale for GPT’s responses, all dyads related to a particular stereotype were presented (in randomized order) within a single chat. For each attitude/stereotype, the same (randomized) order of dyads was presented in prompts to GPT-3.5 and GPT-4. GPT was told that we were interested in what cultural associations *are* as opposed to what they *should be* and instructed to estimate real associations and not to attempt to debias the results. Full materials may be found in the document “GPT as Data Generator Materials 20240227” and transcripts in “GPT as Data Generator Transcripts 20240221”, both at <https://osf.io/sdahr/>.

For analysis, GPT’s responses were treated as analogous to cosine similarity (2). To calculate a measure of relative cultural association – e.g. a greater association of Female with Home and Male with Career, relative to Male with Home and Female with Career – the procedure was followed for calculating the WEAT *D*-score, adapting the method recommended by Charlesworth et

al. (1). The data used in these calculations may be found in the four excel spreadsheets “GPT as Data Generator A-D” at <https://osf.io/sdahr/>.

Although higher-status groups are often preferred to lower-status groups (3), gender attitudes are typically an exception. A wide body of research shows that people tend to implicitly and explicitly associate women more than men to goodness (4-6), and this difference has been replicated on a cultural scale in word embedding research (1). Therefore, we expected responses from GPT to reflect a pattern of greater positivity toward Female words relative to Male words. We expected to see a greater association of Female with Art (and Male with Science) compared to Male with Art (and Female with Science), in line with prior research. Similarly, we expected to see a greater association of Female with Home (and Male with Work) compared to Male with Home (and Female with Work). Finally, we expected to see a greater association of Female with Reading (and Male with Math) compared to Female with Math (and Male with Reading).

Stimuli. The stimuli for this study were taken directly from prior research (1), and with the authors’ permission are reproduced in Table S9.

**Table S9.** *Word Stimuli for GPT as Data Generator Study*

| Category | Word Stimuli                                                                                                           |
|----------|------------------------------------------------------------------------------------------------------------------------|
| Female   | she, her, mommy, mom, girl, mother, lady, sister, mama, momma, sis, grandma, herself                                   |
| Male     | he, his, daddy, dad, boy, father, guy, brother, dada, papa, bro, grandpa, himself                                      |
| Good     | happiness, happy, fun, fantastic, lovable, magical, delight, joy, relaxing, honest, excited, laughter, lover, cheerful |
| Bad      | torture, murder, abuse, wreck, die, disease, disaster, mourning, virus, killer, nightmare, stress, kill, death         |
| Home     | baby, house, home, wedding, kid, family, marry                                                                         |
| Work     | work, office, job, business, trade, activity, act, money                                                               |
| Art      | art, dance, dancing, sing, singing, paint, painting, song, draw, drawing                                               |
| Science  | science, scientist, chemistry, physics, engineer, space, spaceship, astronaut, chemical, microscope                    |
| Reading  | book, read, write, story, word, writing, reading, tale                                                                 |
| Math     | puzzle, number, count, math, counting, calculator, subtraction, addition                                               |

Notes: Ratings were provided by GPT-3.5 and GPT-4 for every possible combination of the words in the “Female” and “Male” (target) categories with those in the other (attribute) categories.

## **Section S7: GPT as Data Generator – Supporting Analyses**

Single-Category WEAT *D*. In the main article, we show that results from both GPT-3.5 and GPT-4 generally aligned with the overall WEAT *D*-scores reported in prior work. In addition, we calculated Single-Category WEAT *D*-scores for each of the concepts, prying apart the degree to which the results in each case were driven by stronger Female-Male associations with one attribute or the other (e.g. with Work vs. Home). Charlesworth et al. (1) report in their Supplemental Materials that in each case, their results appear to be driven jointly by the two attributes. For example, their Work-Home results reflect both a stronger association of Female with Home and of Male with Work. Interestingly, as shown in Table S10, our results diverge from theirs on this count.

In nearly all cases, the overall results in Study 3 are driven primarily by a stronger association of Female with the stereotypically Female category. For example, results from both GPT-3.5 and GPT-4 indicated a stronger association of Female vs. Male with Good; They did not however suggest a parallel stronger association of Male with Bad. In the case of the Home-Work and Reading-Math tasks, Female was associated with *both* categories, but more strongly with the stereotype-consistent category. Only in the case of the Art-Science task was the main result clearly jointly driven by an association of Female with Art and of Male with Science.

**Table S10.** *Single-Category WEAT D-scores from ChatGPT Responses*

|                              | <b>GPT-3.5</b>      | <b>GPT-4</b>        |
|------------------------------|---------------------|---------------------|
|                              |                     |                     |
| Overall pro-Female Attitude  | WEAT $D = 1.00$     | WEAT $D = 0.57$     |
| Female vs. Male, Good        | SC-WEAT $D = 1.15$  | SC-WEAT $D = 0.72$  |
| Female vs. Male, Bad         | SC-WEAT $D = 0.02$  | SC-WEAT $D = -0.01$ |
|                              |                     |                     |
| Overall Female, Art-Science  | WEAT $D = 1.16$     | WEAT $D = 1.46$     |
| Female vs. Male, Art         | SC-WEAT $D = 0.89$  | SC-WEAT $D = 2.26$  |
| Female vs. Male, Science     | SC-WEAT $D = -0.97$ | SC-WEAT $D = -0.92$ |
|                              |                     |                     |
| Overall Female, Home-Work    | WEAT $D = 0.40$     | WEAT $D = 0.45$     |
| Female vs. Male, Home        | SC-WEAT $D = 1.72$  | SC-WEAT $D = 2.15$  |
| Female vs. Male, Work        | SC-WEAT $D = 0.99$  | SC-WEAT $D = 0.52$  |
|                              |                     |                     |
| Overall Female, Reading-Math | WEAT $D = 0.73$     | WEAT $D = 0.96$     |
| Female vs. Male, Reading     | SC-WEAT $D = 2.82$  | SC-WEAT $D = 1.89$  |
| Female vs. Male, Math        | SC-WEAT $D = 0.07$  | SC-WEAT $D = 0.51$  |

Notes: These calculations treat Male/Female as the target categories, and thus differ somewhat from those reported in Table S11.



In addition, we calculated the other form of Single-Category (and bidirectional) WEAT *D*-scores, with the different stereotypically female (versus stereotypically male) attributes (e.g. Good versus Bad) now treated as the target categories, and “Female” and “Male” treated as the attribute categories. The results are presented in Table S11.

It is interesting to note, in Table S11, that GPT-3.5 generated an unexpectedly higher association of Good (vs. Bad) with Male than it did with Female, and also that the same association using data from GPT-4 comes out as only modestly higher for Female than for Male. At first glance, this seems misaligned with the results of the overall WEAT *D* analysis. The reason for this pattern is that this is a *standardized* score that uses the standard deviation of difference scores from within the relevant categories. There was more variance in the difference scores (avg. of Good minus Bad words) for Female words compared to Male words, particularly for GPT-3.5. For example, using data from GPT-3.5, Grandma and Mommy showed notably larger difference scores (2.59, 2.63) than Female and Herself (1.61, 1.33). The corresponding differences were less sharply contrasting for the Male category, with Grandpa and Daddy showing only slightly larger difference scores (1.82, 1.68) than Male and Himself (1.26, 1.51). The result is that while the Female comparisons show larger average differences between Good and Bad words, to calculate the SC-WEAT *D*-scores, they are divided by a larger standard deviation, leading the results to be smaller.

**Table S11.** *Alternative Single-Category WEAT D-scores from ChatGPT Responses*

|                            | <b>GPT-3.5</b>     | <b>GPT-4</b>       |
|----------------------------|--------------------|--------------------|
|                            |                    |                    |
| Overall Good-Bad by Gender | WEAT $D = 1.14$    | WEAT $D = 0.68$    |
| Good vs. Bad, Female       | SC-WEAT $D = 4.67$ | SC-WEAT $D = 3.85$ |
| Good vs. Bad, Male         | SC-WEAT $D = 6.48$ | SC-WEAT $D = 3.43$ |
|                            |                    |                    |
| Art-Science, by Gender     | WEAT $D = 1.36$    | WEAT $D = 1.68$    |
| Art vs. Science, Female    | SC-WEAT $D = 2.58$ | SC-WEAT $D = 4.56$ |
| Art vs. Science, Male      | SC-WEAT $D = 1.10$ | SC-WEAT $D = 2.51$ |
|                            |                    |                    |
| Home-Work, by Gender       | WEAT $D = 0.66$    | WEAT $D = 0.97$    |
| Home vs. Work, Female      | SC-WEAT $D = 1.19$ | SC-WEAT $D = 2.20$ |
| Home vs. Work, Male        | SC-WEAT $D = 0.07$ | SC-WEAT $D = 0.79$ |
|                            |                    |                    |
| Reading-Math, by Gender    | WEAT $D = 0.86$    | WEAT $D = 1.36$    |
| Reading vs. Math, Female   | SC-WEAT $D = 3.25$ | SC-WEAT $D = 6.35$ |
| Reading vs. Math, Male     | SC-WEAT $D = 1.00$ | SC-WEAT $D = 3.76$ |

Notes: These calculations treat Male/Female as the attribute categories, and thus differ somewhat from those reported in Table S10.

## Section S8: GPT as Novel Data Predictor – Detailed Methods

Full Design. As noted in the main manuscript, we asked GPT-3.5 and GPT-4 to predict cross-cultural patterns of implicit and explicit attitudes, drawn from Charlesworth et al. (7). These data, which draw on Implicit Association Tests and explicit questionnaires from over 2.3 million participants across 34 countries, were first published after ChatGPT’s training cutoff at the time of data collection. We selected three attitudes from this study: Sexuality Attitudes, Age Attitudes and Gender Science/Liberal Arts stereotypes. For each, we separately asked GPT-3.5 and GPT-4 to predict the cross-cultural patterns of explicit attitudes (as measured by survey items) and implicit

attitudes (as measured by Implicit Association Tests). (Full prompts may be found in the document “GPT as Novel Data Predictor Materials 20240227” at <https://osf.io/sdahr/>.) Each prompt (Explicit/Implicit x GPT-3.5/4 x 3 attitudinal constructs) was run 5 times, for a total of 60 sets of predictions. In each prompt, we instructed ChatGPT that a more positive D-score (or higher score on explicit items) should indicate a stronger preference in the typical direction from prior research, e.g. a preference for “Straight” over “Gay,” or “Young” over “Old.” We examined each prediction, summarily tracking correlations between GPT’s prediction by country and the actual country-level implicit and explicit attitudes and stereotypes reported in the recent paper (7). Full transcripts from this study may be found in the document “GPT as Novel Data Predictor Transcripts 20240221” at <https://osf.io/sdahr/>.

### **Section S9: GPT as Novel Data Predictor – Supporting Analyses**

Correlations reported here and in the main article were conducted in Stata, and were averaged from correlation tables in Excel. These correlation tables may be found in the spreadsheet “GPT as Data Predictor Correlation Tables 20240228” at <https://osf.io/sdahr/>.

Real Country-Level Implicit-Explicit Correlations. In the main article, we note that when data are aggregated to the country level, the real correlations between Implicit and Explicit attitudes in the Project Implicit International Data Set “vary by attitude object but are generally not strong.” Here, we report these analyses in more detail. In calculating these correlations, we used country-level summary data based on Implicit Association Tests (“implicit” attitudes) and the 7-point Likert items (“explicit” attitudes). Note that the real-world data set also includes a second type of explicit item

(feeling thermometers). We did not ask ChatGPT to predict these feeling thermometers, so we do not examine these items here. Note also that to examine the psychometric properties of the data set, Charlesworth et al. (7) report individual-level correlations between implicit and explicit attitudes. These analyses are distinct from our analyses: we focus on results that have been aggregated by country since these are what we asked ChatGPT to predict. The strongest country-level implicit-explicit correlation was for Sexuality Attitudes:  $r = 0.5332$ ,  $P = 0.0008$ . For Age Attitudes, country-level implicit and explicit results were essentially uncorrelated:  $r = 0.2030$ ,  $P = 0.2350$ . Finally, for country-level Gender Science/Liberal Arts stereotypes, there was a trend toward a *negative* correlation between implicit and explicit results:  $r = -0.3135$ ,  $P = 0.0626$ .

GPT-3.5's difficulty with the task. GPT-3.5 had difficulty completing this task coherently. Though we'd instructed it clearly as to the direction in which results should be scored (i.e. what a positive IAT *D*-score should mean), it was frequently not clear whether it was following this direction. For example, in one chat it would rate a few countries as the *highest* on Implicit Sexuality bias, and then in another it would rank the same countries as the *lowest*. Thus, to ensure we were interpreting the data correctly, across different chats, we unexpectedly needed to ask GPT-3.5 a follow-up question, requesting the direction - i.e. whether a positive average *D*-score reflected a preference for gay or for straight - in which it had scored the countries, even though we had instructed it clearly on this point. Moreover, even after asking this, GPT-3.5 sometimes appeared confused about its own response. In the fourth chat in which it was asked about Implicit Sexuality Attitudes, GPT-3.5's responses were heavily and negatively correlated with those from all other

implicit and explicit predictions (all  $r_s < -0.69$ ,  $P_s < 0.0001$ ), suggesting it misrepresented the interpretation of its own  $D$ -scores. For the main analyses, this response was thus reverse-scored.

This problem was even more pronounced in GPT-3.5's predictions of Age Attitudes and Gender Science/Liberal Arts stereotypes, rendering the data difficult to interpret. For Sexuality Attitudes, we were able to clearly understand what GPT-3.5 was attempting to communicate simply by asking a follow-up question. Given the answers to these questions, as described above, nine of its ten predictions were highly correlated with one another, and one traveled in precisely the opposite direction. We are confident in interpreting this as a mistake on GPT-3.5's part and simply reverse-scoring the answers in this one chat. For the other two attitudes, no such clear pattern arose.

For these latter attitudes, GPT appeared deeply confused about the direction of its own scoring for implicit attitudes, and even at times for explicit attitudes. Indeed, GPT-3.5 became so hopelessly befuddled at times that it avoided answering the follow-up question about the direction of its scoring. For example, when asked what direction the  $D$ -scores in its predictions were scored, instead of answering for its own answers it would apologize and then report the direction in which  $D$ -scores for Implicit Age Attitudes are *typically* scored, necessitating a second follow-up to again clarify the scoring in its own responses. As can be seen in the full correlation tables and transcripts (see "GPT as Data Predictor Correlation Tables 20240228" and "GPT as Novel Data Predictor Transcripts" at <https://osf.io/sdahr/>), some of GPT-3.5's answers on these chats were correlated with each other, but others were

uncorrelated or negatively correlated, with the result that it is impossible to confidently interpret even the intended direction of GPT-3.5's answers. Put more concisely, GPT-3.5's answers for Age Attitudes and Gender Science stereotypes were completely incoherent. An examination of the correlations across chats (scored in the direction GPT-3.5 instructed us) will render this point more clearly.

Age Attitudes. GPT-3.5's five different predictions of Explicit Age Attitudes were, on average, uncorrelated with each other (mean  $r = -0.068$ , SD = 0.548). GPT-3.5's five different predictions of Implicit Age Attitudes were similarly uncorrelated (mean  $r = -0.137$ , SD = 0.509), as were the correlations between its implicit and explicit predictions (mean  $r = 0.021$ , SD = 0.556). Its collective predictions were similarly uncorrelated with real country-level patterns of Explicit (mean  $r = -0.010$ , SD = 0.459) and Implicit (mean  $r = -0.175$ , SD = 0.273) Age Attitudes.

Gender Science/Arts Stereotypes. As with Age Attitudes, GPT-3.5's predictions of Gender Science/Liberal Arts stereotypes seem to suffer from a lack of coherence. Its different explicit stereotype predictions were collectively uncorrelated with each other (mean  $r = -0.115$ , SD = 0.537), as were its different implicit stereotype predictions (mean  $r = -0.186$ , SD = 0.628). Similarly, explicit stereotype predictions were collectively uncorrelated with implicit stereotype predictions (mean  $r = -0.040$ , SD = 0.641). Considering these patterns, it is not surprising that GPT-3.5's predictions were also collectively uncorrelated with real country-level Explicit (mean  $r = -0.009$ , SD = 0.363) and Implicit (mean  $r = -0.044$ , SD = 0.239) Gender stereotypes.

Full Summary Statistics. In the main article, we reported average correlations for most of the main tasks ChatGPT completed in Study 4. For the sake of completeness and ease of reference, we report all average correlations and the standard deviations of these correlations in Table S12.

**Table S12.** *Summary Statistics for GPT as Novel Data Predictor*

| <b>Task/Model</b>     | <b>Correlations:<br/>Explicit with<br/>other<br/>Explicit<br/>Predictions</b> | <b>Correlations:<br/>Implicit with<br/>other<br/>Implicit<br/>Predictions</b> | <b>Correlations:<br/>Implicit with<br/>Explicit<br/>Predictions</b> | <b>Correlations:<br/>All Predictions<br/>with Real<br/>Explicit<br/>Results</b> | <b>Correlations:<br/>All Predictions<br/>with Real<br/>Implicit<br/>Results</b> |
|-----------------------|-------------------------------------------------------------------------------|-------------------------------------------------------------------------------|---------------------------------------------------------------------|---------------------------------------------------------------------------------|---------------------------------------------------------------------------------|
| Sexuality,<br>GPT-3.5 | $M_r = 0.875$<br>$SD_r = 0.082$                                               | $M_r = 0.879$<br>$SD_r = 0.031$                                               | $M_r = 0.778$<br>$SD_r = 0.071$                                     | $M_r = 0.602$<br>$SD_r = 0.073$                                                 | $M_r = -0.014$<br>$SD_r = 0.118$                                                |
| Sexuality,<br>GPT-4   | $M_r = 0.957$<br>$SD_r = 0.026$                                               | $M_r = 0.946$<br>$SD_r = 0.022$                                               | $M_r = 0.952$<br>$SD_r = 0.027$                                     | $M_r = 0.714$<br>$SD_r = 0.040$                                                 | $M_r = 0.152$<br>$SD_r = 0.065$                                                 |
| Age,<br>GPT-3.5       | $M_r = -0.068$<br>$SD_r = 0.548$                                              | $M_r = -0.137$<br>$SD_r = 0.509$                                              | $M_r = 0.021$<br>$SD_r = 0.556$                                     | $M_r = -0.010$<br>$SD_r = 0.459$                                                | $M_r = -0.175$<br>$SD_r = 0.273$                                                |
| Age,<br>GPT-4         | $M_r = 0.645$<br>$SD_r = 0.234$                                               | $M_r = 0.726$<br>$SD_r = 0.161$                                               | $M_r = 0.664$<br>$SD_r = 0.271$                                     | $M_r = -0.395$<br>$SD_r = 0.172$                                                | $M_r = -0.120$<br>$SD_r = 0.196$                                                |
| Gender,<br>GPT-3.5    | $M_r = -0.115$<br>$SD_r = 0.537$                                              | $M_r = -0.186$<br>$SD_r = 0.628$                                              | $M_r = -0.040$<br>$SD_r = 0.641$                                    | $M_r = -0.009$<br>$SD_r = 0.363$                                                | $M_r = -0.044$<br>$SD_r = 0.239$                                                |
| Gender,<br>GPT-4      | $M_r = 0.363$<br>$SD_r = 0.435$                                               | $M_r = 0.868$<br>$SD_r = 0.064$                                               | $M_r = 0.499$<br>$SD_r = 0.425$                                     | $M_r = -0.304$<br>$SD_r = 0.205$                                                | $M_r = -0.007$<br>$SD_r = 0.207$                                                |

Note:  $M_r$  values reflect the mean of each set of correlations, and  $SD_r$  values reflect the standard deviation of the correlations in each group. For example, GPT-3.5 made 5 different predictions of Explicit Sexuality Attitudes. The first cell in this table displays the mean and standard deviation of the 10 different possible pairwise correlations among these 5 predictions.

Supplemental Analysis of Sexuality Attitudes. In the main article, we reported that both GPT-3.5 and especially GPT-4 made implicit predictions that correlated highly with other implicit predictions, and explicit predictions that correlated highly with other explicit predictions, suggesting a degree of reliability in how it approached predicting the cross-country results for Sexuality attitudes. Critically, though, the correlations between implicit and explicit predictions were nearly as high for GPT-3.5 and just as high for GPT-

4, suggesting that ChatGPT did not leverage substantially different information when predicting the (more novel) Implicit Sexuality Attitudes compared to the (less novel) Explicit ones. In general, GPT-3.5 and especially GPT-4 did a relatively strong job of predicting Explicit Sexuality Attitudes (mean  $r$ s = 0.602 and 0.714 respectively), but both did a poor job predicting the more novel implicit results (mean  $r$ s = -0.014 and 0.152). Here, we report some additional analysis of these attitudes.

Exploratory analysis suggested that GPT-3.5 had particular difficulty predicting Sexuality Attitudes in East Asia: It generally predicted these countries would show some of the highest Sexuality bias, when in actuality they showed some of the lowest. Removing China, South Korea, Japan and Taiwan from the sample slightly improves GPT-3.5's prediction of country-level explicit attitudes (mean  $r$  = 0.652, SD = 0.074), and notably improves its prediction of implicit attitudes (mean  $r$  = 0.335, SD = 0.083). Exploratorily removing East Asian countries improves GPT-4's performance similarly to its predecessor's. With this exclusion, average predictions of explicit attitudes improve slightly (mean  $r$  = 0.768, SD = 0.043), and predictions of implicit attitudes improve notably (mean  $r$  = 0.472, SD = 0.049). These findings are consistent with prior reports suggesting limitations in ChatGPT's training data in East Asian languages (8, Supplemental Materials).

Full Correlation Tables may be found in the spreadsheet "GPT as Data Predictor Correlation Tables 20240228" at <https://osf.io/sdahr/>.

Examinations of Media Coverage by Attitude Domain. It is likely that GPT has more access to information about Sexuality Attitudes in its training data, providing a plausible explanation for why it was far more successful in




predicting cross-cultural patterns in these attitudes, particularly when measured explicitly, compared to the others. However, it proved highly difficult to find research directly comparing the media coverage for different issues, to provide evidence for (or against) this argument. Thus, as mentioned in the main article, we conducted three simple but convergent tests to establish whether Sexuality Attitudes are likely to be better represented in worldwide media (and thus GPT's training data), relative to Age Attitudes and Gender Science/Liberal Arts Stereotypes.

First, we conducted a series of searches using Microsoft's Bing search engine, conducted at the time of the manuscript's authorship. Searching (without quotation marks) for "Gay rights" yielded ~128,000,000 results, while searching for "Women's rights" yielded ~58,300,000, "Elder Rights" ~2,780,000, "Elderly Rights" ~2,920,000, and "Old-age rights" ~4,220,000. Searching "Gender Rights" yielded ~170,000,000 results, even more than "Gay Rights". However, it should be noted that this search is far broader than the actual stereotype (Gender-Science) we were studying. We thus conducted a second, more specific set of searches. Searching for "Gender-Science Stereotypes" yielded ~206,000 results; An inexact but related search ("Gender-Science Attitudes" yielded ~355,000; "Age Attitudes" yielded ~164,000. Searches related to Sexuality Attitudes, in contrast, yielded sharply more results. Specifically, searching "Gay Attitudes" yielded ~1,640,000 results, and searching "Homosexuality Attitudes" yielded ~1,490,000 results. These results, while not decisive, suggest that Sexuality Attitudes are better represented in overall worldwide discourse, relative to Age Attitudes and Gender Science/Liberal Arts Stereotypes.


Second, we conducted a series of searches using Google's Ngram Viewer, a tool that allows you to track the frequency with which phrases appear over time within the Google Books corpus. As seen in Figure S1, the phrase "attitudes toward homosexuality" appeared more than "attitudes toward the elderly" between the years 2000 and 2019, and the phrase "gender science stereotypes" appeared too rarely to be plotted. As seen in Figure S2, when we slightly broadened the gender-related search term ("gender in science"), it now appeared in the plot. However, "attitudes toward homosexuality" still consistently appeared far more than the other two phrases. Though this an imperfect test – different searches might yield different results – it nevertheless provides convergent evidence that Sexuality Attitudes have received more media representation in recent years relative to Age Attitudes and Gender Science/Liberal Arts Stereotypes.

**Figure S1.**



Notes: Figure S1 depicts results of a Google Ngram search for "attitudes toward homosexuality", "attitudes toward the elderly" and "gender science stereotypes" between 2000 and 2019.

**Figure S2.**



Notes: Figure S2 depicts results of a Google Ngram search for “attitudes toward homosexuality”, “attitudes toward the elderly” and “gender in science” between 2000 and 2019.

Finally, since our primary interest is in the representation of the different attitudes/stereotypes in ChatGPT’s training data, we decided to probe this question more directly by asking GPT-4 about the relative global coverage of the different issues. Specifically, we asked GPT-4 the following question: “Across cultures, which of these things are spoken of the most? Attitudes toward gays and homosexuality, attitudes towards the elderly, or gender representation in science?” Note that as a more conservative test, we broadened our discussion to be around “gender representation in science” as opposed to a more specific term such as “gender science stereotypes”. As seen in the further supplemental materials, in all 10 chats GPT offered the view that “Attitudes toward gays and homosexuality” have received more global media coverage than the other two topics. (See the document “Chats with GPT-4 about media coverage of different topics 20231222” at <https://osf.io/sdahr/>.)

Taken together, these three results provide convergent evidence that Sexuality Attitudes receive more global coverage than Age Attitudes and Gender-Science Stereotypes. This lends credence to the theory that information access may explain ChatGPT's greater success at predicting Sexuality Attitudes.

Probing GPT-4's knowledge of Implicit & Explicit Stereotypes. At several points in the main article, we claim that GPT likely has within its knowledge base an understanding that Implicit and Explicit attitudes are somewhat independent. This point is important since it makes the result surprising that GPT often seemed to lean on identical information when predicting these two types of attitudes. While confident in our assumption here, to test it, we directly probed GPT-4 about this topic, requesting the following: "What is the difference between explicit and implicit attitudes and how are the two to be measured? Do these kinds of attitudes tend to correspond to one another?"

GPT-4's response to this query can be found in the document "Chat with GPT-4 about implicit and explicit attitudes 20240509" at <https://osf.io/sdahr/>. GPT's answer confirms our contention that the chatbot is aware of the distinction between these types of attitudes. The answer makes some points which might be contended, e.g. unilaterally referring to Implicit Attitudes as "unconscious," a point of scholarly debate. Nevertheless, it clearly demonstrates GPT's knowledge of the two constructs. Most critically, GPT correctly confirms that these attitudes often operate independently, saying for example: "In many cases, particularly with socially sensitive subjects, explicit and implicit attitudes can diverge significantly." It is therefore interesting that GPT displays no meaningful differences in the patterns of

predictions on, for example, Implicit versus Explicit Sexuality Attitudes, a good example of such a socially sensitive subject.

## SI References

1. T. E. S. Charlesworth, V. Yang, T. C. Mann, B. Kurdi, M. R. Banaji, Gender stereotypes in natural language: Word embeddings show robust consistency across child and adult language corpora of more than 65 million words. *Psychol. Sci.* **32**, 218–240 (2021).
2. A. Caliskan, J. J. Bryson, A. Narayanan, Semantics derived automatically from language corpora contain human-like biases. *Science* **356**, 183–186 (2017).
3. J. T. Jost, M. R. Banaji, B. A. Nosek, A decade of system justification theory: Accumulated evidence of conscious and unconscious bolstering of the status quo. *Polit. Psychol.* **25**, 881–919 (2004).
4. A. H. Eagly, A. Mladinic, Are people prejudiced against women? Some answers from research on attitudes, gender stereotypes, and judgments of competence. *Eur. Rev. Soc. Psychol.* **5**, 1–35 (1994).
5. B. A. Nosek, M. R. Banaji, The go/no-go association task. *Soc. Cognit.* **19**, 625–666 (2001).
6. L. A. Rudman, S. A. Goodwin, Gender differences in automatic in-group bias: Why do women like women more than men like men? *J. Pers. Soc. Psychol.* **87**, 494–509 (2004).
7. T. E. S. Charlesworth, M. Navon, Y. Rabinovich, N. Lofaro, B. Kurdi, The project implicit international dataset: Measuring implicit and explicit social group attitudes and stereotypes across 34 countries (2009–2019). *Behav. Res. Methods* **55**, 1413–1440 (2023).
8. T. Brown *et al.*, Language models are few-shot learners. *Adv. Neural Inf. Process. Syst.* **33**, 1877–1901 (2020).